


Evaluation of the Context-Free Path Querying Algorithm Based on Matrix Multiplication

Nikita Mishin, Iaroslav Sokolov, Egor Spirin, Vladimir Kutuev, Egor
Nemchinov, Sergey Gorbatyuk, **Semyon Grigorev**

JetBrains Research, Programming Languages and Tools Lab
Saint Petersburg University

June 30, 2019

Context-Free Path Querying



Navigation through a graph

- Are nodes A and B on the same level of hierarchy?
- Is there a path of form $Up^n Down^n$?
- Find all paths of form $Up^n Down^n$ which start from the node A

Context-Free Path Querying: Relational Query Semantics

- $\mathbb{G} = (\Sigma, N, P)$ — context-free grammar in normal form
 - ▶ $A \rightarrow BC$, where $A, B, C \in N$
 - ▶ $A \rightarrow x$, where $A \in N, x \in \Sigma$
 - ▶ $L(\mathbb{G}, A) = \{\omega \mid A \Rightarrow^* \omega\}$
- $G = (V, E, L)$ — directed graph
 - ▶ $v \xrightarrow{l} u \in E$
 - ▶ $L \subseteq \Sigma$
- $\omega(\pi) = \omega(v_0 \xrightarrow{l_0} v_1 \xrightarrow{l_1} \dots \xrightarrow{l_{n-2}} v_{n-1} \xrightarrow{l_{n-1}} v_n) = l_0 l_1 \dots l_{n-1}$
- $R_A = \{(n, m) \mid \exists n\pi m, \text{ such that } \omega(\pi) \in L(\mathbb{G}, A)\}$

Matrix-Based Algorithm

Algorithm Context-Free Path Querying by Matrix Multiplication

```
1: function CONTEXTFREEPATHQUERYING( $D, G$ )
2:    $n \leftarrow$  the number of nodes in  $D$ 
3:    $E \leftarrow$  the directed edge-relation from  $D$ 
4:    $P \leftarrow$  the set of production rules in  $G$ 
5:    $T \leftarrow$  the matrix  $n \times n$  in which each element is  $\emptyset$ 
6:   for all  $(i, x, j) \in E$  do ▷ Matrix initialization
7:      $T_{i,j} \leftarrow T_{i,j} \cup \{A \mid (A \rightarrow x) \in P\}$ 
8:   while matrix  $T$  is changing do
9:      $T \leftarrow T \cup (T \times T)$  ▷ Transitive closure  $T^{cf}$  calculation
10:  return  $T$ 
```

Matrix-Based Algorithm: Technical Details

- T can be represented as set of Boolean matrices: one matrix per nonterminal
- The algorithm can be implemented in terms of Boolean matrices multiplication
- All matrices can be statically allocated in memory

Research Questions

- Does GPGPUs utilization for CFPQ improve performance in comparison with the CPU version?

Research Questions

- Does GPGPUs utilization for CFPQ improve performance in comparison with the CPU version?
- Is it possible to achieve higher performance by using existing libraries for operations over matrices or do we need to create our own specialized solution?

Research Questions

- Does GPGPUs utilization for CFPQ improve performance in comparison with the CPU version?
- Is it possible to achieve higher performance by using existing libraries for operations over matrices or do we need to create our own specialized solution?
- Can we achieve high performance with high-level languages?

Research Questions

- Does GPGPUs utilization for CFPQ improve performance in comparison with the CPU version?
- Is it possible to achieve higher performance by using existing libraries for operations over matrices or do we need to create our own specialized solution?
- Can we achieve high performance with high-level languages?
- Can we improve performance with sparse matrix representation?

[Scipy] Sparse matrices multiplication by using **Scipy** in **Python**

CPU-Based Implementations

[Scipy] Sparse matrices multiplication by using **Scipy** in **Python**

[M4RI] Dense matrices multiplication by using **m4ri** library which implements the Method of Four Russians in **C**

[GPU4R] Our own implementation of the Method of Four Russians
in **CUDA C**

[GPU4R] Our own implementation of the Method of Four Russians
in **CUDA C**

[GPU_N] Our own implementation of the naïve boolean matrix
multiplication in **CUDA C**

GPGPU-Based Implementations

[GPU4R] Our own implementation of the Method of Four Russians in **CUDA C**

[GPU_N] Our own implementation of the naïve boolean matrix multiplication in **CUDA C**

[GPU_Py] Our own implementation of naïve boolean matrix multiplication in **Python** by using **numba** compiler

- [CuSprs]
- ▶ Rustam Azimov, 2018, “Context-free Path Querying by Matrix Multiplication”
 - ▶ Implementation is based on NVIDIA cuSPARSE library (**CUDA C, GPGPU**)

Reference Implementations

[CuSprs]

- ▶ Rustam Azimov, 2018, “Context-free Path Querying by Matrix Multiplication”
- ▶ Implementation is based on NVIDIA cuSPARSE library (**CUDA C, GPGPU**)

[CYK]

- ▶ X. Zhang et al, 2016, “Context-free path queries on RDF graphs”
- ▶ CYK-based algorithm implemented in **Java** (CPU)

[RDF]

- ▶ The set of the real-world RDF files (ontologies)

- ▶ Queries:

$$G_4 : s \rightarrow SCOR\ s\ SCO \mid TR\ s\ T \mid SCOR\ SCO \mid TR\ T$$
$$G_5 : s \rightarrow SCOR\ s\ SCO \mid SCO$$

Dataset

[RDF]

- ▶ The set of the real-world RDF files (ontologies)


- ▶ Queries:

$G_4 : s \rightarrow SCOR\ s\ SCO \mid TR\ s\ T \mid SCOR\ SCO \mid TR\ T$

$G_5 : s \rightarrow SCOR\ s\ SCO \mid SCO$

[Worst]

- ▶ The input graph is two cycles of coprime lengths with one shared vertex



- ▶ Query: $G_1 : s \rightarrow A\ s\ B \mid A\ B$

- [Full]
- ▶ The input graph is sparse, but the result is a full graph
 - ▶ Queries:
 $G_2 : s \rightarrow s \ s \mid A$
 $G_3 : s \rightarrow s \ s \ s \mid A$

[Full]

- ▶ The input graph is sparse, but the result is a full graph

- ▶ Queries:

$$G_2 : s \rightarrow s \ s \mid A$$

$$G_3 : s \rightarrow s \ s \ s \mid A$$

[Sparse]

- ▶ Sparse graphs are generated by GTgraph

- ▶ Query: $G_1 : s \rightarrow A \ s \ B \mid A \ B$

- OS: Ubuntu 18.04
- CPU: Intel core i7 8700k 3,7HGz
- RAM: DDR4 32 Gb
- GPGPU: Geforce 1080Ti (11Gb RAM)

Evaluation: [RDF]²

| RDF | | | Query G_4 | | | | | | |
|-------------|-----|------|-------------|------|-------|-------|--------|--------|------------------|
| Name | #V | #E | Scipy | M4RI | GPU4R | GPU_N | GPU_Py | CuSprs | CYK ¹ |
| atm-prim | 291 | 685 | 3 | 2 | 2 | 1 | 5 | 269 | 515285 |
| biomed | 341 | 711 | 3 | 5 | 2 | 1 | 5 | 283 | 420604 |
| foaf | 256 | 815 | 2 | 9 | 2 | < 1 | 5 | 270 | 5027 |
| funding | 778 | 1480 | 4 | 7 | 4 | 1 | 5 | 279 | 499 |
| generations | 129 | 351 | 3 | 3 | 2 | < 1 | 5 | 273 | 6091 |
| people_pets | 337 | 834 | 3 | 3 | 3 | 1 | 7 | 284 | 82081 |
| pizza | 671 | 2604 | 6 | 8 | 3 | 1 | 6 | 292 | 3233587 |
| skos | 144 | 323 | 2 | 4 | 2 | < 1 | 5 | 273 | 1044 |
| travel | 131 | 397 | 3 | 5 | 2 | < 1 | 6 | 268 | 13971 |
| unv-bnch | 179 | 413 | 2 | 4 | 2 | < 1 | 5 | 266 | 20981 |
| wine | 733 | 2450 | 7 | 6 | 4 | 1 | 7 | 294 | 4075319 |

¹Results from X. Zhang et al, 2016, "Context-Free Path Queries on RDF Graphs"

²Time in milliseconds

Evaluation: [Worst]³

| #V | Scipy | M4RI | GPU4R | GPU_N | GPU_Py | CuSprs |
|------|----------|---------|--------|---------|---------|----------|
| 16 | 0.032 | < 0.001 | 0.008 | 0.002 | 0.027 | 0.309 |
| 32 | 0.118 | 0.001 | 0.034 | 0.008 | 0.136 | 0.441 |
| 64 | 0.476 | 0.041 | 0.133 | 0.032 | 0.524 | 0.988 |
| 128 | 2.194 | 0.226 | 0.562 | 0.129 | 2.751 | 3.470 |
| 256 | 15.299 | 1.994 | 3.088 | 0.544 | 11.883 | 15.317 |
| 512 | 121.287 | 23.204 | 13.685 | 2.499 | 43.563 | 102.269 |
| 1024 | 1593.284 | 528.521 | 88.064 | 19.357 | 217.326 | 1122.055 |
| 2048 | - | - | - | 325.174 | - | - |

³Time in seconds

Evaluation: [Worst]³

| #V | Scipy | M4RI | GPU4R | GPU_N | GPU_Py | CuSprs |
|------|----------|---------|--------|---------|---------|----------|
| 16 | 0.032 | < 0.001 | 0.008 | 0.002 | 0.027 | 0.309 |
| 32 | 0.118 | 0.001 | 0.034 | 0.008 | 0.136 | 0.441 |
| 64 | 0.476 | 0.041 | 0.133 | 0.032 | 0.524 | 0.988 |
| 128 | 2.194 | 0.226 | 0.562 | 0.129 | 2.751 | 3.470 |
| 256 | 15.299 | 1.994 | 3.088 | 0.544 | 11.883 | 15.317 |
| 512 | 121.287 | 23.204 | 13.685 | 2.499 | 43.563 | 102.269 |
| 1024 | 1593.284 | 528.521 | 88.064 | 19.357 | 217.326 | 1122.055 |
| 2048 | - | - | - | 325.174 | - | - |

³Time in seconds

Evaluation: [Sparse]⁴

| Graph | Scipy | M4RI | GPU4R | GPU_N | GPU_Py | CuSprs |
|------------|---------|----------|--------|--------|--------|---------|
| G5k-0.001 | 10.352 | 0.647 | 0.113 | 0.041 | 0.216 | 5.729 |
| G10k-0.001 | 37.286 | 2.395 | 0.435 | 0.215 | 1.331 | 35.937 |
| G10k-0.01 | 97.607 | 1.455 | 0.273 | 0.138 | 0.763 | 47.525 |
| G10k-0.1 | 601.182 | 1.050 | 0.223 | 0.114 | 0.859 | 395.393 |
| G20k-0.001 | 150.774 | 11.025 | 1.842 | 1.274 | 6.180 | - |
| G40k-0.001 | - | 97.841 | 11.663 | 8.393 | 37.821 | - |
| G80k-0.001 | - | 1142.959 | 88.366 | 65.886 | - | - |

⁴Time in seconds

Evaluation: [Sparse]⁴

| Graph | Scipy | M4RI | GPU4R | GPU_N | GPU_Py | CuSprs |
|------------|---------|----------|--------|--------|--------|---------|
| G5k-0.001 | 10.352 | 0.647 | 0.113 | 0.041 | 0.216 | 5.729 |
| G10k-0.001 | 37.286 | 2.395 | 0.435 | 0.215 | 1.331 | 35.937 |
| G10k-0.01 | 97.607 | 1.455 | 0.273 | 0.138 | 0.763 | 47.525 |
| G10k-0.1 | 601.182 | 1.050 | 0.223 | 0.114 | 0.859 | 395.393 |
| G20k-0.001 | 150.774 | 11.025 | 1.842 | 1.274 | 6.180 | - |
| G40k-0.001 | - | 97.841 | 11.663 | 8.393 | 37.821 | - |
| G80k-0.001 | - | 1142.959 | 88.366 | 65.886 | - | - |

⁴Time in seconds

Evaluation: [Sparse]⁴

| Graph | Scipy | M4RI | GPU4R | GPU_N | GPU_Py | CuSprs |
|------------|---------|----------|--------|--------|--------|---------|
| G5k-0.001 | 10.352 | 0.647 | 0.113 | 0.041 | 0.216 | 5.729 |
| G10k-0.001 | 37.286 | 2.395 | 0.435 | 0.215 | 1.331 | 35.937 |
| G10k-0.01 | 97.607 | 1.455 | 0.273 | 0.138 | 0.763 | 47.525 |
| G10k-0.1 | 601.182 | 1.050 | 0.223 | 0.114 | 0.859 | 395.393 |
| G20k-0.001 | 150.774 | 11.025 | 1.842 | 1.274 | 6.180 | - |
| G40k-0.001 | - | 97.841 | 11.663 | 8.393 | 37.821 | - |
| G80k-0.001 | - | 1142.959 | 88.366 | 65.886 | - | - |

⁴Time in seconds

Evaluation: [Full]⁵

| #V | Query G ₂ | | | | | |
|-------|----------------------|----------|---------|---------|---------|----------|
| | Scipy | M4RI | GPU4R | GPU_N | GPU_Py | CuSprs |
| 100 | 0.007 | 0.002 | 0.002 | < 0.001 | 0.003 | 0.278 |
| 200 | 0.040 | 0.003 | 0.002 | 0.001 | 0.004 | 0.279 |
| 500 | 0.480 | 0.003 | 0.003 | 0.001 | 0.004 | 0.329 |
| 1000 | 3.741 | 0.007 | 0.005 | 0.001 | 0.006 | 0.571 |
| 2000 | 40.309 | 0.063 | 0.019 | 0.003 | 0.017 | 1.949 |
| 5000 | 651.343 | 0.366 | 0.125 | 0.038 | 0.150 | 99.651 |
| 10000 | - | 1.932 | 0.552 | 0.315 | 0.840 | 1029.042 |
| 25000 | - | 33.236 | 7.252 | 5.314 | 15.521 | - |
| 50000 | - | 360.035 | 58.751 | 44.611 | 129.641 | - |
| 80000 | - | 1292.817 | 256.579 | 190.343 | 641.260 | - |

⁵Time in seconds

Conclusion

- GPGPUs utilization significantly increases the performance of CFPQ

Conclusion

- GPGPUs utilization significantly increases the performance of CFPQ
- High performance libraries utilization is a good idea
 - ▶ But not always: M4RI (CPU) is better then cuSPARSE (GPGPU)

Conclusion

- GPGPUs utilization significantly increases the performance of CFPQ
- High performance libraries utilization is a good idea
 - ▶ But not always: M4RI (CPU) is better then cuSPARSE (GPGPU)
- Automatic translation from a high-level language to GPGPU language provides a good balance between performance and implementation complexity

Conclusion

- GPGPUs utilization significantly increases the performance of CFPQ
- High performance libraries utilization is a good idea
 - ▶ But not always: M4RI (CPU) is better then cuSPARSE (GPGPU)
- Automatic translation from a high-level language to GPGPU language provides a good balance between performance and implementation complexity
- Sparse matrix representation is important for performance

Conclusion

- GPGPUs utilization significantly increases the performance of CFPQ
- High performance libraries utilization is a good idea
 - ▶ But not always: M4RI (CPU) is better then cuSPARSE (GPGPU)
- Automatic translation from a high-level language to GPGPU language provides a good balance between performance and implementation complexity
- Sparse matrix representation is important for performance
- Dataset is published: both graphs and queries
- Implementations are available on GitHub
- Link: <https://github.com/SokolovYaroslav/CFPQ-on-GPGPU>

- Investigate implemented algorithms to explain nontrivial behaviors
- Create open extensible platform for CFPQ algorithms comparison
- Evaluate other CFPQ algorithms
 - ▶ Sparse matrices
 - ▶ Distributed matrix multiplication
 - ▶ LL- and LR-based algorithms
- Add new data and queries to the dataset
 - ▶ Bigger RDFs
 - ▶ Static code analysis

Contact Information

- Semyon Grigorev:
 - ▶ s.v.grigoriev@spbu.ru
 - ▶ Semen.Grigorev@jetbrains.com
- Nikita Mishin: mishinnikitam@gmail.com
- Iaroslav Sokolov: sokolov.yas@gmail.com
- Egor Spirin: egor@spirin.tech
- Vladimir Kutuev: vladimir.kutuev@gmail.com
- Egor Nemchinov: nemchegor@gmail.com
- Sergey Gorbatyuk: sergeygorbatyuk171@gmail.com

- Dataset and algorithm implementations:
<https://github.com/SokolovYaroslav/CFPQ-on-GPGPU>

Thanks!