

POSTER: Optimizing GPU Programs by Partial Evaluation

Aleksey Tyurin
Saint Petersburg State University
alekseytyurinspb@gmail.com

Daniil Berezun
JetBrains Research
daniil.berezun@jetbrains.com

Semyon Grigorev*
Saint Petersburg State University
s.v.grigoriev@spbu.ru

Abstract

While GPU utilization allows one to speed up computations to the orders of magnitude, memory management remains the bottleneck making it often a challenge to achieve the desired performance. Hence, different memory optimizations are leveraged to make memory being used more effectively. We propose an approach automating memory management utilizing partial evaluation, a program transformation technique that enables data accesses to be pre-computed, optimized, and embedded into the code, saving memory transactions. An empirical evaluation of our approach shows that the transformed program could be up to 8 times as efficient as the original one in the case of CUDA C naïve string pattern matching algorithm implementation.

CCS Concepts • Software and its engineering → General programming languages; Source code generation;

Keywords GPU, CUDA, Partial Evaluation

1 Introduction

Performance of GPU-based solutions critically depends on memory management, since most applications tend to be bandwidth bound, requiring the data to be allocated and accessed most effectively. Thereby, memory optimizations appear to be in a prevailing significance and addressed in a huge amount of research.

To address issues related to the lack of available GPU memory, sophisticated memory pooling techniques are used leveraging memory swapping and sharing [9]. Further, since GPU memory access latency varies between various memory types, techniques like shared memory register spilling or automatic shared memory allocation are utilized to save global memory transactions and reduce cache contention [6, 8]. However, memory type could be used inefficiently due to

improper type-specific access patterns that could aggravate memory access latency in virtue of entailing memory transaction number increase. We propose an approach that inherently constitutes a runtime memory optimization that reduces the overall memory transaction number thereby addressing the mentioned issues. The approach is based on the typical common workflow feature of GPU-based solutions described below.

Suppose one has created an interactive solution for huge data analysis allowing the end-user to sequentially write GPU kernel processed queries to a dataset. Let a query be relatively small and data be large, resulting in execution time being significant. Simple examples of such scenarios are multiple patterns matching, database querying, and convolutional filtering. The GPU kernel, in this case, should be generic and have at least two parameters: a query and data. Being created by the user the query can be used as static, i.e. already known and constant, data for the kernel runtime optimization since it remains unchanged during the host code execution. Specifically, the kernel could be *partially evaluated* at runtime with respect to the query being issued.

Partial evaluation or *program specialization* is a well-known optimization technique that given a program and part of its input data, called *static*, specializes the program with respect to the data, producing another, optimized, program which if given only the remaining part of input data, called *dynamic*, yields the same results as the original program would have produced being executed given both parts of the input data [1, 2].

Application of specialization for one of the described scenarios, namely database querying, has been recently known to significantly improve CPU-based queries execution performance [7]. It appears that partial evaluation is able to achieve similar results for GPU-based applications, optimizing memory accesses. Particularly, considering the problem of multiple patterns matching, where the patterns are *static* during the execution, the memory access result for a particular pattern could be evaluated at specialization time and be embedded as a value into the code at compilation time, rather than being compiled to a load instruction for some specific memory type. More precisely, partial evaluation lets the values to be accessed through instructions cache, avoiding issuing load transactions. Thus, in the paper, we propose to apply the partial evaluation technique for GPU code optimization and evaluate it on one of the described scenarios.

* Also with JetBrains Research.

Permission to make digital or hard copies of part or all of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for third-party components of this work must be honored. For all other uses, contact the owner/author(s).

PPoPP '20, February 22–26, 2020, San Diego, CA, USA

© 2020 Copyright held by the owner/author(s).

ACM ISBN 978-1-4503-6818-6/20/02.

<https://doi.org/10.1145/3332466.3374507>

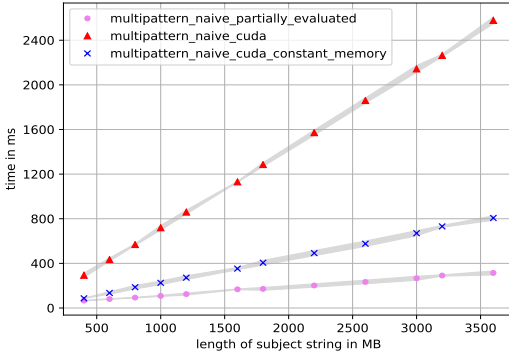


Figure 1. Multiple string pattern matching evaluation

2 Evaluation

The approach has been evaluated considering the problem of *file carving* [5] that stands for extracting files from raw data in a field of *cyber forensics*. To extract a file, a specific file header should be detected, and the set of relatively short file headers becomes static at the search query execution time. Hence, the query could be specialized with respect to the headers.

Basically, we compare¹ naïve multiple string matching algorithm implementations: one leveraging partial evaluator developed as part of AnyDSL framework [4], and two baseline ones in CUDA C with global and constant memory for header access respectively. All the implementations invoke the algorithm in a separate thread for each position in the subject string. The algorithm simply iterates over the headers char-array searching for a match, if it encounters a mismatch, it jumps to the next header forward through the array. During specialization the partial evaluator performs loop unrolling and evaluates unrolled memory access instructions for the runtime-known headers. The results² of the evaluation³ are presented in Figure 1.

Since mismatches happen quite often inducing thread divergence, such access pattern hurts coalescing, increasing the overall number of memory transactions. Given that, the performance speedup of a partially evaluated algorithm achieved on a raw 4 Gb data piece is up to 8 compared to CUDA C version with global memory and up to 3 with constant one as illustrated in Figure 1. Namely, the partially evaluated version spends about 300 ms for searching while global and constant memory CUDA C versions making it in 800 ms and 2500 ms respectively. Specialization overhead constitutes 2 sec, which makes performance improvement significant in case of analysing 1 Tb storage.

¹The approach has been evaluated on Ubuntu 18.04 system with *Intel Core i7-6700* processor, 8Gb of RAM and *Pascal-based GeForce GTX 1070* GPU with 8Gb device memory.

²The points are the average kernel running time and the grey regions are the area of standard deviation.

³The subject string has been taken from a hard drive and patterns to be searched have been taken from a taxonomy of file headers specifications [3].

3 Conclusion

In this work, we apply partial evaluation to optimize GPU programs, showing that this optimization technique can significantly improve performance, not requiring manual manipulation with source code and implementation of additional memory management routines.

The partial evaluator being used assumes the programs to be written with special *DSL*. Nevertheless, partial evaluation could be applied to general-purpose languages, and CUDA C, and the upcoming research is dedicated to the generalization of the technique so as the partial evaluation could be applied at runtime, during GPU-based application execution. Moreover, the performance of partial evaluator should be improved in order to decrease specialization overhead.

Finally, specialization could produce code with a huge number of variables, making the compiler to spill excessive registers. A workaround would be a pipelining of specialization with other optimization techniques, e.g. with shared memory register spilling [6].

Acknowledgments

This material is based upon work supported by the National Science Foundation under Grant No. nnnnnnnn and Grant No. mmmmmmm. Any opinions, findings, and conclusions or recommendations expressed in this material are those of the author and do not necessarily reflect the views of the National Science Foundation.

References

- [1] Neil D. Jones. 1996. An Introduction to Partial Evaluation. *ACM Comput. Surv.* 28, 3 (1996), 480–503. <https://doi.org/10.1145/243439.243447>
- [2] Neil D. Jones, Carsten K. Gomard, and Peter Sestoft. 1993. *Partial Evaluation and Automatic Program Generation*. Prentice-Hall, Inc., Upper Saddle River, NJ, USA.
- [3] Gary Kessler. 2019. GCK’S FILE SIGNATURES TABLE. https://www.garykessler.net/library/file_sigs.html. Accessed: 2019-10-31.
- [4] Roland Leissa, Klaas Boesche, Sebastian Hack, Arsène Pérard-Gayot, Richard Membarth, Philipp Slusallek, André Müller, and Bertil Schmidt. 2018. AnyDSL: A Partial Evaluation Framework for Programming High-performance Libraries. *Proc. ACM Program. Lang.* 2, OOPSLA, Article 119 (Oct. 2018), 30 pages. <https://doi.org/10.1145/3276489>
- [5] Digambar Povar and V. K. Bhadrar. 2011. Forensic Data Carving. In *Digital Forensics and Cyber Crime*, Ibrahim Baggili (Ed.). Springer Berlin Heidelberg, Berlin, Heidelberg, 137–148.
- [6] Putt Sakdhnagool, Amit Sabne, and Rudolf Eigenmann. 2019. RegDem: Increasing GPU Performance via Shared Memory Register Spilling. *ArXiv abs/1907.02894* (2019).
- [7] Eugene Sharygin, Ruben Buchatskiy, Roman Zhuykov, and Arseny Sher. 2018. Runtime Specialization of PostgreSQL Query Executor. In *Perspectives of System Informatics*, Alexander K. Petrenko and Andrei Voronkov (Eds.). Springer International Publishing, Cham, 375–386.
- [8] Xinfeng Xie, Jason Cong, and Yun Liang. 2018. ICCAD : U : Optimizing GPU Shared Memory Allocation in Automated Cto-CUDA Compilation.
- [9] Junzhe Zhang, Sai Ho Yeung, Yao Shu, Bingsheng He, and Wei Wang. 2019. Efficient Memory Management for GPU-based Deep Learning Systems. *arXiv:cs.DC/1903.06631*