

POSTER: Optimizing GPU Programs By Partial Evaluation

Aleksey Tyurin, Daniil Berezun, **Semyon Grigorev**

JetBrains Research, Programming Languages and Tools Lab
Saint Petersburg University


February 24, 2020

GPGPU Architecture


Navigation through a graph

- Global memory

😊 Big




GPGPU Architecture



Navigation through a graph

- Global memory
 - ☺ Big
 - ☹ Slow
- Shared memory
 - ☺ Fast

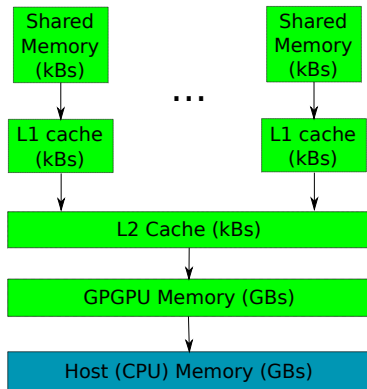
GPGPU Architecture



Navigation through a graph

- Global memory
 - 😊 Big
 - 😞 Slow
- Shared memory
 - 😊 Fast
 - 😞 Relatively small
 - 😞 Manual allocation management
- Constant memory
 - 😊 Fast

GPGPU Architecture



Navigation through a graph

- Global memory
 - 😊 Big
 - 😞 Slow
- Shared memory
 - 😊 Fast
 - 😞 Relatively small
 - 😞 Manual allocation management
- Constant memory
 - 😊 Fast
 - 😞 Only for appropriate access pattern
 - 😞 Small
 - 😞 Static allocation
- Memory traffic is a bottleneck

Data Processing

- Substring matching
- Filtering by using Hidden Markov Models (HMM)

- Substring matching
- Filtering by using Hidden Markov Models (HMM)

```
__global__ void estimateSimilarity
                (int* filterParams, int* data, ...)
{
    __shared__ int cachedFilterParams[size];

    /*some code to load filterParams
       to cachedFilterParams*/
    ...
}
```

Data Processing

- Substring matching
- Filtering by using Hidden Markov Models (HMM)

```
__global__ void estimateSimilarity
    (int* filterParams, int* data, ...)
{
    __shared__ int cachedFilterParams[size];

    /*some code to load filterParams
       to cachedFilterParams*/
    ...
}
```


Data Processing

- Substring matching
- Filtering by using Hidden Markov Models (HMM)

```
__global__ void estimateSimilarity
    (int* filterParams, int* data, ...)
{
    __shared__ int cachedFilterParams[size];

    /*some code to load filterParams
       to cachedFilterParams*/
    ...
}
```

Data Processing

- Substring matching
- Filtering by using Hidden Markov Models (HMM)

```
__global__ void estimateSimilarity
                (int* filterParams, int* data, ...)
{
    __shared__ int cachedFilterParams[size];

    /*some code to load filterParams
       to cachedFilterParams*/

    ...
}
```

Big Data Processing

- Substring matching \Rightarrow Data curving (digital forensics)
- Filtering by using Hidden Markov Models (HMM) \Rightarrow Homology search (bioinformatics)

Big Data Processing

- Substring matching \Rightarrow Data curving (digital forensics)
- Filtering by using Hidden Markov Models (HMM) \Rightarrow Homology search (bioinformatics)

```
__global__ void estimateSimilarity
                (int* filterParams, int* data, ...)
{
    ...
}
```

Big Data Processing

- Substring matching \Rightarrow Data curving (digital forensics)
- Filtering by using Hidden Markov Models (HMM) (bioinformatics)

Many data chunks
 \Rightarrow many runs
of procedure

```
__global__ void estimateSimilarity  
                (int* filterParams, int* data, ...)  
{  
    ...  
}
```

Big Data Processing

- Substring matching \Rightarrow Data curving (digital forensics)
- Filtering by using Hidden Markov Models (HMM) (bioinformatics)

One filter for
many data chunks

Many data chunks
 \Rightarrow many runs
of procedure

```
__global__ void estimateSimilarity  
    (int* filterParams, int* data, ...)  
{  
    ...  
}
```

Big Data Processing

- Substring matching \Rightarrow Data curving (digital forensics)
- Filtering by using Hidden Markov Models (HMM) (bioinformatics)

One filter for
many data chunks

Many data chunks
 \Rightarrow many runs
of procedure

```
__global__ void estimateSimilarity  
    (int* filterParams, int* data, ...)  
{  
    ...  
}
```

filterParams is a static during one data porcessing session.

Big Data Processing

- Substring matching \Rightarrow Data curving (digital forensics)
- Filtering by using Hidden Markov Models (HMM) (bioinformatics)

One filter for
many data chunks

Many data chunks
 \Rightarrow many runs
of procedure

```
__global__ void estimateSimilarity  
    (int* filterParams, int* data, ...)  
{  
    ...  
}
```

filterParams is a static during one data porcessing session.
How can we use this fact to optimize our procedure?

Partial Evaluation



!!! Framework

[Scipy] Sparse matrices multiplication by using **Scipy** in **Python**

!!! Framework

[Scipy] Sparse matrices multiplication by using **Scipy** in **Python**

[M4RI] Dense matrices multiplication by using **m4ri** library which implements the Method of Four Russians in **C**

[GPU4R] Our own implementation of the Method of Four Russians
in **CUDA C**

[GPU4R] Our own implementation of the Method of Four Russians
in **CUDA C**

[GPU_N] Our own implementation of the naïve boolean matrix
multiplication in **CUDA C**

Evaluation: Data Curving

[GPU4R] Our own implementation of the Method of Four Russians in **CUDA C**

[GPU_N] Our own implementation of the naïve boolean matrix multiplication in **CUDA C**

[GPU_Py] Our own implementation of naïve boolean matrix multiplication in **Python** by using **numba** compiler

[CuSprs]

- ▶ Rustam Azimov, 2018, "Context-free Path Querying by Matrix Multiplication"
- ▶ Implementation is based on NVIDIA cuSPARSE library (**CUDA C, GPGPU**)

Evaluation: !!!

[CuSprs]

- ▶ Rustam Azimov, 2018, “Context-free Path Querying by Matrix Multiplication”
- ▶ Implementation is based on NVIDIA cuSPARSE library (**CUDA C, GPGPU**)

[CYK]

- ▶ X. Zhang et al, 2016, “Context-free path queries on RDF graphs”
- ▶ CYK-based algorithm implemented in **Java** (CPU)

Limitations

[RDF]

- ▶ The set of the real-world RDF files (ontologies)

- ▶ Queries:

$$G_4 : s \rightarrow SCOR\ s\ SCO \mid TR\ s\ T \mid SCOR\ SCO \mid TR\ T$$

$$G_5 : s \rightarrow SCOR\ s\ SCO \mid SCO$$

Limitations

[RDF]

- ▶ The set of the real-world RDF files (ontologies)


- ▶ Queries:

$$G_4 : s \rightarrow SCOR \ s \ SCO \mid TR \ s \ T \mid SCOR \ SCO \mid TR \ T$$

$$G_5 : s \rightarrow SCOR \ s \ SCO \mid SCO$$

[Worst]

- ▶ The input graph is two cycles of coprime lengths with one shared vertex



- ▶ Query: $G_1 : s \rightarrow A \ s \ B \mid A \ B$

- [Full]
- ▶ The input graph is sparse, but the result is a full graph
 - ▶ Queries:
 $G_2 : s \rightarrow s \ s \mid A$
 $G_3 : s \rightarrow s \ s \ s \mid A$

[Full]

- ▶ The input graph is sparse, but the result is a full graph

- ▶ Queries:

$$G_2 : s \rightarrow s \ s \mid A$$

$$G_3 : s \rightarrow s \ s \ s \mid A$$

[Sparse]

- ▶ Sparse graphs are generated by GTgraph

- ▶ Query: $G_1 : s \rightarrow A \ s \ B \mid A \ B$

Conclusion

- OS: Ubuntu 18.04
- CPU: Intel core i7 8700k 3,7GHz
- RAM: DDR4 32 Gb
- GPGPU: NVIDIA GeForce 1080Ti (11Gb RAM)

- Investigate implemented algorithms to explain nontrivial behaviors
- Create open extensible platform for CFPQ algorithms comparison
- Evaluate other CFPQ algorithms
 - ▶ Sparse matrices
 - ▶ Distributed matrix multiplication
 - ▶ LL- and LR-based algorithms
- Add new data and queries to the dataset
 - ▶ Bigger RDFs
 - ▶ Static code analysis

Contact Information

- Semyon Grigorev:
 - ▶ s.v.grigoriev@spbu.ru
 - ▶ Semen.Grigorev@jetbrains.com
- Nikita Mishin: mishinnikitam@gmail.com
- Iaroslav Sokolov: sokolov.yas@gmail.com
- Egor Spirin: egor@spirin.tech
- Vladimir Kutuev: vladimir.kutuev@gmail.com
- Egor Nemchinov: nemchegor@gmail.com
- Sergey Gorbatyuk: sergeygorbatyuk171@gmail.com

- Dataset and algorithm implementations:
<https://github.com/SokolovYaroslav/CFPQ-on-GPGPU>

Thanks!