

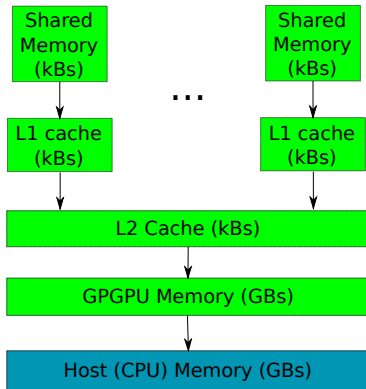
POSTER: Optimizing GPU Programs By Partial Evaluation

Aleksey Tyurin, Daniil Berezun, **Semyon Grigorev**

JetBrains Research, Programming Languages and Tools Lab
Saint Petersburg University

February 24, 2020

GPGPU memory hierarchy

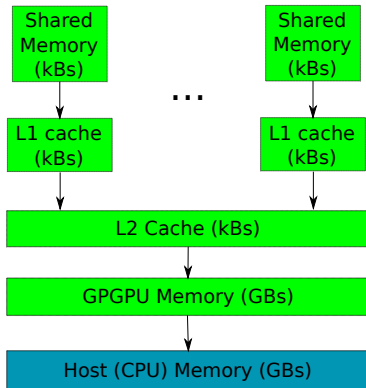


GPGPU Architecture

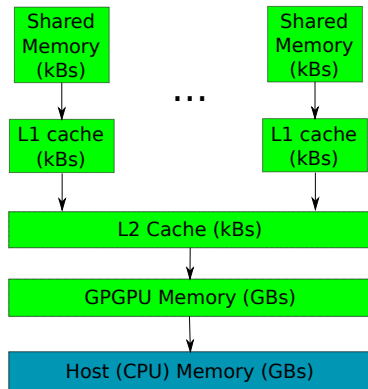
GPGPU memory hierarchy

- Global memory

- 😊 Big
- 😞 Slow



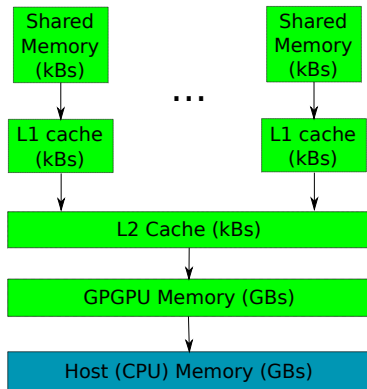
GPGPU Architecture



GPGPU memory hierarchy

- Global memory
 - 😊 Big
 - 😞 Slow
- Shared memory
 - 😊 Fast
 - 😞 Relatively small
 - 😞 Manual allocation management

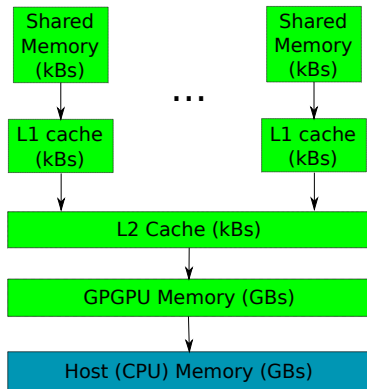
GPGPU Architecture



GPGPU memory hierarchy

- Global memory
 - 😊 Big
 - ☹️ Slow
- Shared memory
 - 😊 Fast
 - ☹️ Relatively small
 - ☹️ Manual allocation management
- Constant memory
 - 😊 Fast
 - ☹️ Only for appropriate access pattern
 - ☹️ Small
 - ☹️ Static allocation

GPGPU Architecture



GPGPU memory hierarchy

- Global memory
 - 😊 Big
 - 😞 Slow
- Shared memory
 - 😊 Fast
 - 😞 Relatively small
 - 😞 Manual allocation management
- Constant memory
 - 😊 Fast
 - 😞 Only for appropriate access pattern
 - 😞 Small
 - 😞 Static allocation
- Memory traffic is a bottleneck

- Substring matching
- Filtering by using Hidden Markov Models (HMM)

Data Processing

- Substring matching
- Filtering by using Hidden Markov Models (HMM)

```
__global__ void handleData
    (int* filterParams, int* data, ...)
{
    __shared__ int cachedFilterParams[size];

    /*some code to load filterParams
       to cachedFilterParams*/
    ...
}
```


Data Processing

- Substring matching
- Filtering by using Hidden Markov Models (HMM)

```
__global__ void handleData
    (int* filterParams, int* data, ...)
{
    __shared__ int cachedFilterParams[size];

    /*some code to load filterParams
       to cachedFilterParams*/
    ...
}
```

Data Processing

- Substring matching
- Filtering by using Hidden Markov Models (HMM)

```
__global__ void handleData
    (int* filterParams, int* data, ...)
{
    __shared__ int cachedFilterParams[size];

    /*some code to load filterParams
       to cachedFilterParams*/
    ...
}
```

Data Processing

- Substring matching
- Filtering by using Hidden Markov Models (HMM)

```
__global__ void handleData
    (int* filterParams, int* data, ...)
{
    __shared__ int cachedFilterParams[size];

    /*some code to load filterParams
       to cachedFilterParams*/

    ...
}
```

Big Data Processing

- Substring matching \Rightarrow Data curving (cyber forensics)
- Filtering by using Hidden Markov Models (HMM) \Rightarrow Homology search (bioinformatics)

Big Data Processing

- Substring matching \Rightarrow Data curving (cyber forensics)
- Filtering by using Hidden Markov Models (HMM) \Rightarrow Homology search (bioinformatics)

```
__global__ void handleData
                (int* filterParams, int* data, ...)
{
    ...
}
```

Big Data Processing

- Substring matching \Rightarrow Data curving (cyber forensics)
- Filtering by using Hidden Markov Models (HMM) (bioinformatics)

Many data chunks
 \Rightarrow many runs
of procedure

```
__global__ void handleData  
                (int* filterParams, int* data, ...)  
{  
    ...  
}
```

Big Data Processing

- Substring matching \Rightarrow Data curving (cyber forensics)
- Filtering by using Hidden Markov Models (HMM) (bioinformatics)

One filter for
many data chunks

Many data chunks
 \Rightarrow many runs
of procedure

```
__global__ void handleData  
    (int* filterParams, int* data, ...)  
{  
    ...  
}
```

Big Data Processing

- Substring matching \Rightarrow Data curving (cyber forensics)
- Filtering by using Hidden Markov Models (HMM) (bioinformatics)

One filter for
many data chunks

Many data chunks
 \Rightarrow many runs
of procedure

```
__global__ void handleData  
    (int* filterParams, int* data, ...)  
{  
    ...  
}
```

`filterParams` is a static during one data porcessing session.

Big Data Processing

- Substring matching \Rightarrow Data curving (cyber forensics)
- Filtering by using Hidden Markov Models (HMM) (bioinformatics)

One filter for
many data chunks

Many data chunks
 \Rightarrow many runs
of procedure

```
__global__ void handleData  
    (int* filterParams, int* data, ...)  
{  
    ...  
}
```

`filterParams` is a static during one data porcessing session.

How can we use this fact to optimize our procedure?

Partial Evaluation or Specialization

$$\underbrace{\llbracket \text{handleData} \rrbracket}_{\text{handleData}}[\text{filterParams}, \text{data}] = \underbrace{\overbrace{\llbracket \text{mix} \rrbracket[\text{handleData}, \text{filterParams}]]^{\text{partial evaluator}}}_{\text{handleData}_{\text{mix}}}[\text{data}]$$

Partial Evaluation or Specialization

$$\underbrace{\llbracket \text{handleData} \rrbracket}_{\text{handleData}}[\text{filterParams}, \text{data}] = \overbrace{\llbracket \text{mix} \rrbracket[\llbracket \text{handleData}, \text{filterParams} \rrbracket]}^{\text{partial evaluator}}[\text{data}]$$

$\underbrace{\hspace{15em}}_{\text{handleData}_{\text{mix}}}$

```
handleData (filterParams, data)
{
    res = new List()
    for d in data
        for e in filterParams
            if d % e == 0
                then res.Add(d)
    return res
}
```

Partial Evaluation or Specialization

$$\underbrace{\llbracket \text{handleData} \rrbracket}_{\text{handleData}}[\text{filterParams}, \text{data}] = \overbrace{\llbracket \text{mix} \rrbracket[\llbracket \text{handleData}, \text{filterParams} \rrbracket]}^{\text{partial evaluator}}[\text{data}]$$

$\text{handleData}_{\text{mix}}$

$\llbracket \llbracket \text{mix} \rrbracket[\text{handleData}, [2; 3]] \rrbracket$

```
handleData (filterParams, data)
{
    res = new List()
    for d in data
        for e in filterParams
            if d % e == 0
                then res.Add(d)
    return res
}
```

Partial Evaluation or Specialization

$$\underbrace{\llbracket \text{handleData} \rrbracket}_{\text{handleData}}[\text{filterParams}, \text{data}] = \overbrace{\llbracket \text{mix} \rrbracket[\llbracket \text{handleData}, \text{filterParams} \rrbracket]}^{\text{partial evaluator}}[\text{data}]$$

$\text{handleData}_{\text{mix}}$

$\llbracket \text{mix} \rrbracket[\text{handleData}, [2; 3]]$

```
handleData (filterParams, data)
{
    res = new List()
    for d in data
        for e in filterParams
            if d % e == 0
                then res.Add(d)
    return res
}
```

```
handleData (data)
{
    res = new List()
    for d in data
        if d % 2 == 0 ||
           d % 3 == 0
            then res.Add(d)
    return res
}
```

Partial Evaluation or Specialization

$$\underbrace{\llbracket \text{handleData} \rrbracket}_{\text{handleData}}[\text{filterParams}, \text{data}] = \overbrace{\llbracket \text{mix} \rrbracket[\llbracket \text{handleData}, \text{filterParams} \rrbracket]}^{\text{partial evaluator}}[\text{data}]$$

$\text{handleData}_{\text{mix}}$

$\llbracket \llbracket \text{mix} \rrbracket[\text{handleData}, [2; 3]] \rrbracket$

```
handleData (filterParams, data)
{
    res = new List()
    for d in data
        for e in filterParams
            if d % e == 0
                then res.Add(d)
    return res
}
```

```
handleData (data)
{
    res = new List()
    for d in data
        if d % 2 == 0 ||
           d % 3 == 0
            then res.Add(d)
    return res
}
```

- AnyDSL framework for specialization
 - ▶ Special DSL which can be specialized and comiled
 - ▶ Ahead-of-time specialization

Evaluation Setup

- AnyDSL framework for specialization
 - ▶ Special DSL which can be specialized and compiled
 - ▶ Ahead-of-time specialization
- Algorithms
 - ▶ Naïve multiple substring matching
 - ▶

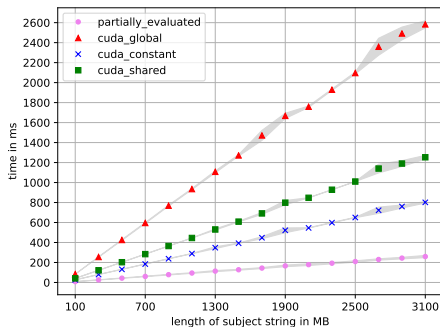
Evaluation Setup

- AnyDSL framework for specialization
 - ▶ Special DSL which can be specialized and compiled
 - ▶ Ahead-of-time specialization
- Algorithms
 - ▶ Naïve multiple substring matching
 - ▶
- Environment

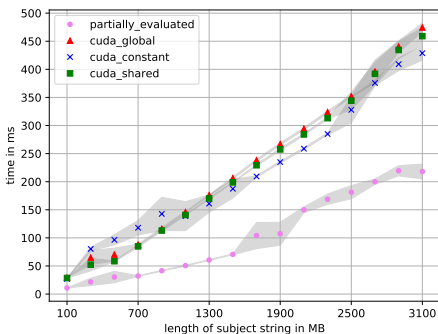
GTX-1070 Environment
T4

Evaluation: Substring Matching

- Application: data curving
- Subject string: byte sequence from real hard drive
- Patterns: 16 file signatures from GCK's file signatures table¹



Results for GTX-1070




Results for T4


¹https://www.garykessler.net/library/file_sigs.html

Evaluation: 2D Convolution

- Application: image processing
- Subject image: random image (16384 * 16384) 1Gb size
- Filters: random square filters with diameter 13 to 255



Results for GTX-1070



Results for T4

Conclusion

- Partial evaluation improves performance of GPGPU procedures
 - ▶ !!!
 - ▶ !!!

- Switch to CUDA C partial evaluator
 - ▶ LLVM.mix: partial evaluator for LLVM IR

Future Research

- Switch to CUDA C partial evaluator
 - ▶ LLVM.mix: partial evaluator for LLVM IR
- Reduce specialization overhead
 - ▶ To be applicable in run-time

- Switch to CUDA C partial evaluator
 - ▶ LLVM.mix: partial evaluator for LLVM IR
- Reduce specialization overhead
 - ▶ To be applicable in run-time
- Integrete with shared memory register spilling
 - ▶ “RegDem: Increasing GPU Performance via Shared Memory Register Spilling” (Putt Sakdhnagool et.al. 2019)

Future Research

- Switch to CUDA C partial evaluator
 - ▶ LLVM.mix: partial evaluator for LLVM IR
- Reduce specialization overhead
 - ▶ To be applicable in run-time
- Integrete with shared memory register spilling
 - ▶ “RegDem: Increasing GPU Performance via Shared Memory Register Spilling” (Putt Sakdhnagool et.al. 2019)
- Evaluate on real-world examples
 - ▶ Homology search in bioinformatics
 - ▶ Graph processing
 - ▶ Graph database querying

Contact Information

- Semyon Grigorev:
 - ▶ s.v.grigoriev@spbu.ru
 - ▶ Semen.Grigorev@jetbrains.com
- Aleksey Tyurin: alekseytyurinspb@gmail.com
- Daniil Berezun: daniil.berezun@jetbrains.com
- Dataset and algorithm implementations:
<https://github.com/SokolovYaroslav/CFPQ-on-GPGPU>

Thanks!