

F# and GPGPU

Semyon Grigorev
St. Petersburg State University
7/9 Universitetskaya nab.
St. Petersburg, 199034, Russia
semen.grigorev@jetbrains.com

Kirill Smirenko
St. Petersburg State University
7/9 Universitetskaya nab.
St. Petersburg, 199034, Russia
k.smirenko@gmail.com

ABSTRACT

Aaabstraaact!!!!

CCS Concepts

•Software and its engineering → Automated static analysis; Software maintenance tools; •Theory of computation → Program analysis; Parsing;

Keywords

GPGPU, OpenCL, F#, metaprogramming, !!!!!

1. INTRODUCTION

GPGPU is popular technique for....

Tools are low level

OpenCL, CUDA etc.

Complex problems, heterogenous platforms: multicore, multi GPGPU etc. Special tools, libs required for development simplification.

F# primitives

General requirements: highlevel language, existing code/dlls/other stuff reusing

Brahma.FSharp – the best platform for GPGPU programming!!!!

2. F# PROGRAMMING LANGUAGE

In this section F# [3] — is a functional-first multiparadigm programming language for .NET platform.

Main important features described.

2.1 Code quotation

2.2 Type providers

2.3 Async MBP etc

3. RELATED WORK

Existing solution for GPGPU programming...

3.1 FSCL

Status [1]

3.2 Alea CUDA

CUDA only

3.3 Managed Cuda etc

Hm...

4. BRAHMA.FSHARP

In this section we present our platform for GPGPU programming in F#.

Blah-Blah-Blah!!!!

4.1 Architecture

Based on F# code quotation to OpenCL translator.

Driver is OpenCL.NET ¹

Picture.

Details of some blocks are described below.

4.2 Translator

4.3 OpenCL specific operations

Atomic functions.

Memory transferring.

4.4 OpenCL type provider

5. EVALUATION

Matrix multiplication

Substring matching

Substring matching with agents [2]

Results of performance test of GPGPU calculation using Brahma.FSharp and MailboxProcessor composition are presented. Problem to solve is substring matching for data carving. Rabin-Karp algorithm was implemented using Brahma.FSharp for substring matching. F# MailboxProcessor used for composing of data reading, data processing on GPGPU, and data processing on CPU. Library for fast and flexible configuration of MailboxProcessors was created. Set of templates for search was fixed. Tests were performed for HDD and SSD storages. Low level sequential reading was implemented. First 16.5 Mb was processed.

¹OpenCL.NET — low level .NET bindings for OpenCL. Project site [visited: 20.06.2017]: <https://openclnet.codeplex.com/>.

- OS: Microsoft Windows 8.1 Pro
- System Type: x64-based PC
- Processor: Intel(R) Core(TM) i7-4790 CPU @ 3.60GHz, 3601 Mhz, 4 Core(s), 8 Logical Processor(s)
- RAM: 16.0 GB
- HDD for test:
 - Model: ST3250410AS
 - Size: 232.88 GB
 - 7200 rpm
- SSD for test
 - Model: INTEL SSDSC2BW240A4
 - Size: 223.57 GB
 - Max read speed: 540 Mb/sec
- GPGPU:
 - NVIDIA GeForce GTX 560 Ti
 - CUDA Cores: 384
 - Core clock: 822 MHz
 - Shader clock: 1645 MHz
 - Memory data rate: 4008 MHz
 - Memory interface: 256-bit
 - Memory bandwidth: 128.26 GB/s
 - Total available graphics memory: 4095 MB
 - Dedicated video memory: 2048 MB GDDR5
 - Shared system memory: 2047 MB

Tables below present results of tests. “buffers for data” — a number of arrays to fill by disc reader for each MailboxProcessor which communicate with GPGPU. “threads” — a number of MailboxProcessors which communicate with GPGPU. In current configuration we have only one GPGU, so all MailboxProcessors use it. For multi-GPGPU systems we can configure k MailboxProcessors for each GPGPU.

In each cell — total time and GPGPU loading graph.

Conclusion: Data reading bufferization can sufficiently increase performance. Especially for HDD, where speed of reading is low. For SSD processing with multi-GPGPU systems may be useful. Data reading is not so critical as for HDD and more than one GPGPU can be fully loaded by using flexible MailboxProcessors configuration. Configuration with two MailboxProcessors and two buffers for each of them can fully load one GPGPU.

6. CONCLUSION

Education.

Graph parsing

Heterogeneous programming generalization. Hopac is better than MBP.

Acknowledgments

We are grateful to the !!!! and !!! for their careful reading, pointing out some mistakes, and invaluable suggestions. This work is supported by grant from JetBrains Research, and by grant UMNIK!!!!.

7. REFERENCES

- [1] G. Coco. *Homogeneous programming, scheduling and execution on heterogeneous platforms*. PhD thesis, Ph.D. Thesis.\Gabriel Coco.-University of Pisa, 2014.-254 p, 2014.
- [2] D. Mikhaylov. *Transparent usage of massively parallel architectures for local optimizations of .NET applications*. PhD thesis, Graduation Thesis.\Dmitry Mikhaylov.-St. Petersburg State University, 2013.-47 p, 2013.
- [3] D. Syme, A. Granicz, and A. Cisternino. *Expert F# 3.0*. Springer, 2012.