

Context-Free Path Querying Can be Fast if Cooked Properly

Arseniy Terekhov

!!!

Saint Petersburg State University
St. Petersburg, Russia

Artyom Khoroshev

!!!

St. Petersburg, Russia

Semyon Grigorev

s.v.grigoriev@spbu.ru

semen.grigorev@jetbrains.com

Saint Petersburg State University
St. Petersburg, Russia
JetBrains Research
St. Petersburg, Russia

ABSTRACT

Recently proposed matrix multiplication based algorithm for context-free path querying (CFPQ) offloads the most performance-critical parts onto boolean matrices multiplication. Thus, it is possible to achieve high performance of CFPQ by means of modern parallel hardware and software. In this paper, we provide results of empirical performance comparison of different implementations of this algorithm on both real-world data and synthetic data for the worst cases.

KEYWORDS

Context-free path querying, transitive closure, graph databases, linear algebra, context-free grammar, GPGPU, CUDA, boolean matrix, matrix multiplication

1 INTRODUCTION

Language-constrained path querying [3], and particularly Context-Free Path Querying (CFPQ) [14], allows one to use formal grammars as constraints for paths: concatenation of the labels along the path is treated as a word, and a constraint on the path is a specification of the language which should contain specific words. CFPQ is widely used for graph-structured data analysis in such domains as biological data analysis, RDF, network analysis. Huge amount of the real-world data makes performance of CFPQ critical for practical tasks. Several algorithms for CFPQ based on such parsing techniques as (G)LL, (G)LR, and CYK are proposed recently [5, 7, 9, 11, 12, 15].

One of the most promising algorithms is a matrix-based algorithm, proposed by Rustam Azimov [2]. This algorithm offloads the most critical computations onto boolean matrices multiplication. As a result, it is easy to implement, and allows one to utilize modern massive-parallel hardware for CFPQ. The implementation provided by the authors utilizes GPGPU by using cuSPARSE¹ library which is a floating point sparse matrices multiplication library. Although it does not use advanced algorithms for boolean matrices, it outperforms existing algorithms.

ww [8] — all is bad, performance is poor, not applicable.

It is necessary to investigate the effect of the specific algorithms and implementation techniques on the performance of CFPQ. One problem is that no publicly available standard dataset for CFPQ algorithms evaluation which includes both graph-structured data and queries exists.

¹cuSparse is a library for sparse matrices multiplication on GPGPU. Official documentation: <https://docs.nvidia.com/cuda/cusparse/index.html>. Access date: 12.03.2019

In this work, we do an empirical performance comparison of several implementations of matrices multiplication based algorithm for CFPQ on both real-world data and synthetic data for the worst cases. We make the following contributions in this paper.

- (1) We provide a number of implementations of the matrix multiplication based CFPQ algorithm, which utilizes different modern software and hardware. The source code is available on GitHub: !!!.
- (2)
- (3) We provide evaluation which shows that !!!.

2 MATRIX-BASED ALGORITHM FOR CFPQ

Matrix-based algorithm for CFPQ was proposed by Rustam Azimov [2]. This algorithm can be expressed in terms of operations over matrices (see listing 1), and it is a sufficient advantage for implementation. It was shown that the utilization of GPGPU improves the context-free path querying performance significantly in comparison to other implementations [2] even if float matrices are used instead of boolean matrices.

Listing 1 Context-free path querying algorithm

```
1: function CONTEXTFREEPATHQUERYING(D, G)
2:    $n \leftarrow$  the number of nodes in  $D$ 
3:    $E \leftarrow$  the directed edge-relation from  $D$ 
4:    $P \leftarrow$  the set of production rules in  $G$ 
5:    $T \leftarrow$  the matrix  $n \times n$  in which each element is  $\emptyset$ 
6:   for all  $(i, x, j) \in E$  do ▷ Matrix initialization
7:      $T_{i,j} \leftarrow T_{i,j} \cup \{A \mid (A \rightarrow x) \in P\}$ 
8:   while matrix  $T$  is changing do
9:      $T \leftarrow T \cup (T \times T)$  ▷ Transitive closure calculation
10:  return  $T$ 
```

Here $D = (V, E)$ is the input graph and $G = (N, \Sigma, P)$ is the input grammar. Each cell of the matrix T contains the set of nonterminals such that $N_k \in T[i, j] \iff \exists p = v_i \dots v_j$ —path in D , such that $N_k \xrightarrow{*}_G \omega(p)$, where $\omega(p)$ is a word formed by the labels along the path p . Thus, this algorithm solves reachability problem, or, according to Hellings [6], processes CFPQs by using relational query semantics.

The performance-critical part of the algorithm is matrix multiplication. Note, that the set of nonterminals is finite, and we can represent the matrix T as a set of boolean matrices: one for each nonterminal. In this case the operation of matrix update is $T_{N_i} \leftarrow T_{N_i} + (T_{N_j} \times T_{N_k})$ for each production $N_i \rightarrow N_j N_k$ in P . Thus we can reduce CFPQ to boolean matrices multiplication. After such transformation, we can apply the next optimization: we can skip update if the matrices T_{N_j} and T_{N_k} have not been changed at the previous iteration.

Thus, the most important part is the efficient implementation of operations over boolean matrices. In this paper, we compare the effects of different approaches to matrices multiplication. All our implementations are based on the optimized version of the algorithm.

3 IMPLEMENTATION

Previous works show [2, 10] that existing linear algebra libraries utilization is a right way to get high-performance CFPQ implementation in minimal effort. But non of these works do not provide evaluation with data storage, only pure time of algorithm execution was measured.

We provide a number of implementations of matrix-based CFPQ algorithm. All of them are based on RedisGraph — we use RedisGraph as a storage and implement CFPQ as an extension by using provided mechanism.

CPU-based implementation. Details on CPU implementation

GPGPU-based implementation. Details on GPGPU implementation

4 DATASET DESCRIPTION

We created and published a dataset for CFPQ algorithms evaluation. This dataset contains both the real-world data and synthetic data for different specific cases, such as the theoretical worst case, or the worst cases specific to matrices representations.

Our goal is to evaluate querying algorithms, not graph storages or graph databases, so all data is presented in a text-based format to simplify usage in different environments. Grammars are in Chomsky Normal Form, and graphs are represented as a list of triples (edges). Some details of the data representation can be found in the Appendix.

The variants of the *same generation query* [1] are an important example of queries that are context-free but not regular, so we use this type of queries in our evaluation. The dataset includes data for the following cases. Each case is a pair of a set of graphs and a set of grammars: each query (grammar) should be applied to each graph.

[RDF] The set of the real-world RDF files (ontologies) from [15] and two variants of the same generation query which describes hierarchy analysis. The first query is the grammar G_4 :

$$\begin{aligned} s &\rightarrow SCOR\ s\ SCO & s &\rightarrow TR\ s\ T \\ s &\rightarrow SCOR\ SCO & s &\rightarrow TR\ T \end{aligned}$$

The second one is the grammar G_5 : $s \rightarrow SCOR\ s\ SCO \mid SCO$.

[Worst] The theoretical worst case for CFPQ time complexity proposed by Hellings [7]: the graph is two cycles of coprime lengths with a single common vertex. The first cycle is labeled by the open bracket and the second cycle is labeled by the close bracket. Query is a grammar for the $A^n B^n$ language. The example of such graph and grammar is presented in figure 1.

[Full] The case when the input graph is sparse, but the result is a full graph. Such a case may be hard for sparse matrices representations. As an input graph, we use a cycle, all edges of which are labeled by the same token. As a query we use two grammars which describe the sequence of tokens of arbitrary length: the simple ambiguous grammar G_2 : $s \rightarrow s\ s \mid A$, and the highly ambiguous grammar G_3 : $s \rightarrow s\ s\ s \mid A$.

[Sparse] Sparse graphs from [4] are generated by the GTgraph graph generator, and emulate realistic sparse data. Names of these graphs have the form G_n-p , where n corresponds to the total number of vertices, and p is the probability that some pair

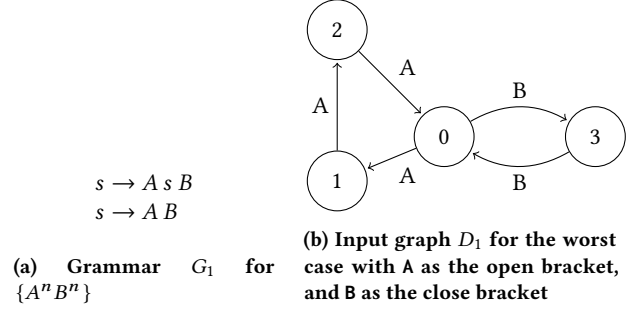


Figure 1: Graph and grammar for the worst case

of vertices is connected. The query is the same generation query represented by the grammar G_1 (figure 1).

5 EVALUATION

We evaluate all the described implementations on all the datasets and the queries presented. We compare our implementations with [2]. We exclude the time required to load data from files. The time required for data transfer is included.

For evaluation, we use a PC with Ubuntu 18.04 installed. It has Intel core i7 8700k 3,7Hz CPU, DDR4 32 Gb RAM, and Geforce 1080Ti GPGPU with 11Gb RAM.

The results of the evaluation are summarized in the tables below. Time is measured in seconds unless specified otherwise. The result for each algorithm is averaged over 10 runs. The cell is left blank if the time limit is exceeded, or if there is not enough memory to allocate the data.

The results of the first dataset **[RDF]** are presented in table 1. We can see, that in this case the running time of all our implementations is smaller than of the reference implementation, and all implementations but **[CuSprs]** demonstrate similar performance. It is obvious that performance improvement in comparison with the first implementation is huge and it is necessary to extend the dataset with new RDFs of the significantly bigger size.

Table 2: Evaluation results for the worst case

#V	Scipy	M4RI	GPU4R	GPU_N	GPU_Py	CuSprs
16	0.032	< 1	0.008	0.002	0.027	0.309
32	0.118	0.001	0.034	0.008	0.136	0.441
64	0.476	0.041	0.133	0.032	0.524	0.988
128	2.194	0.226	0.562	0.129	2.751	3.470
256	15.299	1.994	3.088	0.544	11.883	15.317
512	121.287	23.204	13.685	2.499	43.563	102.269
1024	1593.284	528.521	88.064	19.357	217.326	1122.055
2048	-	-	-	325.174	-	-

Results of the theoretical worst case (**[Worst]** dataset) are presented in table 2. This case is really hard to process: even for a graph of 1024 vertices, the query evaluation time is greater than 10 seconds even for the most performant implementation. We can see, that the running time grows too fast with the number of vertices.

Table 1: RDFs querying results (time in milliseconds)

RDF			Query G_4						Query G_5					
Name	#V	#E	Scipy	M4RI	GPU4R	GPU_N	GPU_Py	CuSprs	Scipy	M4RI	GPU4R	GPU_N	GPU_Py	CuSprs
atm-prim	291	685	3	2	2	1	5	269	1	< 1	1	< 1	2	267
biomed	341	711	3	5	2	1	5	283	4	< 1	1	< 1	5	280
foaf	256	815	2	9	2	< 1	5	270	1	< 1	1	< 1	2	263
funding	778	1480	4	7	4	1	5	279	2	< 1	3	< 1	4	274
generations	129	351	3	3	2	< 1	5	273	1	< 1	1	< 1	2	263
people_pets	337	834	3	3	3	1	7	284	1	< 1	1	< 1	3	277
pizza	671	2604	6	8	3	1	6	292	2	< 1	2	< 1	5	278
skos	144	323	2	4	2	< 1	5	273	< 1	< 1	1	< 1	2	265
travel	131	397	3	5	2	< 1	6	268	1	< 1	1	< 1	3	271
unv-bnch	179	413	2	4	2	< 1	5	266	1	< 1	1	< 1	3	266
wine	733	2450	7	6	4	1	7	294	1	< 1	3	< 1	3	281

Table 3: Sparse graphs querying results

Graph	Scipy	M4RI	GPU4R	GPU_N	GPU_Py	CuSprs
G5k-0.001	10.352	0.647	0.113	0.041	0.216	5.729
G10k-0.001	37.286	2.395	0.435	0.215	1.331	35.937
G10k-0.01	97.607	1.455	0.273	0.138	0.763	47.525
G10k-0.1	601.182	1.050	0.223	0.114	0.859	395.393
G20k-0.001	150.774	11.025	1.842	1.274	6.180	-
G40k-0.001	-	97.841	11.663	8.393	37.821	-
G80k-0.001	-	1142.959	88.366	65.886	-	-

The next is the **[Sparse]** dataset presented in table 3. The evaluation shows that sparsity of graphs (value of parameter p) is important both for implementations which use sparse matrices and for implementations which use dense matrices. Note that the behavior of the sparse matrices based implementation is as expected, but for dense matrices we can see, that more sparse graphs are processed faster. Reasons for such behavior demand further investigation. Note that we estimate only the query execution time, so it is hard to compare our results with the results presented in [4]. Nevertheless, the running time of our **[GPU_N]** implementation is significantly smaller than the one provided in [4].

The last dataset is **[Full]**, and results are shown in table 4. As we expect, this case is very hard for sparse matrices based implementations: the running time grows too fast. This dataset also demonstrates the impact of the grammar size. Both queries specify the same constraints, but the grammar G_3 in CNF contains 2 times more rules then the grammar G_2 , so, the running time for big graphs differs by more than twice.

Finally, we can conclude that GPGPU utilization for CFPQ can significantly improve performance, but more research on advanced optimization techniques should be done. On the other hand, the high-level implementation (**[GPU_Py]**) is comparable with other GPGPU-based implementations. So, it may be a balance between implementation complexity and performance. Highly optimized existing libraries can be of some use: the implementation based on m4ri is faster than the reference implementation and the other CPU-based implementation. Moreover, it is comparable with some GPGPU-based implementations in some cases. Sparse matrices utilization demands more thorough investigation. The main question is if we can create an efficient implementation for sparse boolean matrices multiplication.

6 DISCUSSION

Zeros in sparse matrices.

Overhead on matrices conversion and transferring when run on GPGPU.

No conversion when run on CPU. Better than Neo4j and other BD with non-matrix representation of graphs.

7 CONCLUSION AND FUTURE WORK

Full integration: OpenCupher extension ². LibCypher ³.

GraphBlast [13] – for GPGPU. Multi-GPU graphBlast?

Paths semantics.

Dynamically updated data.

Another direction is the dataset improvement. First of all, it is necessary to collect more data, and more grammars/queries. It is of most importance to add more real-world graphs and more real-world queries to the dataset. Secondly, it is necessary to discuss and fix the data format to be able to evaluate different algorithms. We believe that it is necessary to create a public dataset for CFPQ algorithms evaluation, and collaboration with the community is required.

ACKNOWLEDGMENTS

The research was supported by the Russian Science Foundation grant 18-11-00100 and a grant from JetBrains Research.

²!!!

³!!!

Table 4: Full querying results

#V	Query G_2						Query G_3					
	Scipy	M4RI	GPU4R	GPU_N	GPU_Py	CuSprs	Scipy	M4RI	GPU4R	GPU_N	GPU_Py	CuSprs
100	0.007	0.002	0.002	< 1	0.003	0.278	0.023	0.076	0.005	0.001	0.007	0.290
200	0.040	0.003	0.002	0.001	0.004	0.279	0.105	0.098	0.004	0.001	0.007	0.296
500	0.480	0.003	0.003	0.001	0.004	0.329	1.636	0.094	0.007	0.001	0.010	0.382
1000	3.741	0.007	0.005	0.001	0.006	0.571	13.071	0.106	0.009	0.001	0.009	0.839
2000	40.309	0.063	0.019	0.003	0.017	1.949	93.676	0.108	0.030	0.005	0.026	3.740
5000	651.343	0.366	0.125	0.038	0.150	99.651	1205.421	0.851	0.195	0.075	0.239	201.151
10000	-	1.932	0.552	0.315	0.840	1029.042	-	4.690	1.055	0.648	1.838	-
25000	-	33.236	7.252	5.314	15.521	-	-	70.823	15.240	10.961	36.495	-
50000	-	360.035	58.751	44.611	129.641	-	-	775.765	130.203	91.579	226.834	-
80000	-	1292.817	256.579	190.343	641.260	-	-	-	531.694	376.691	-	-

REFERENCES

- [1] Serge Abiteboul, Richard Hull, and Victor Vianu. 1995. Foundations of Databases.
- [2] Rustam Azimov and Semyon Grigorev. 2018. Context-free Path Querying by Matrix Multiplication. In *Proceedings of the 1st ACM SIGMOD Joint International Workshop on Graph Data Management Experiences & Systems (GRADES) and Network Data Analytics (NDA) (GRADES-NDA '18)*. ACM, New York, NY, USA, Article 5, 10 pages. <https://doi.org/10.1145/3210259.3210264>
- [3] Chris Barrett, Riko Jacob, and Madhav Marathe. 2000. Formal-language-constrained path problems. *SIAM J. Comput.* 30, 3 (2000), 809–837.
- [4] Zhiwei Fan, Jianqiao Zhu, Zuyu Zhang, Aws Albarghouthi, Paraschos Koutris, and Jignesh Patel. 2018. Scaling-Up In-Memory Datalog Processing: Observations and Techniques. *arXiv preprint arXiv:1812.03975* (2018).
- [5] Semyon Grigorev and Anastasiya Ragozina. 2017. Context-free Path Querying with Structural Representation of Result. In *Proceedings of the 13th Central & Eastern European Software Engineering Conference in Russia (CEE-SECR '17)*. ACM, New York, NY, USA, Article 10, 7 pages. <https://doi.org/10.1145/3166094.3166104>
- [6] Jelle Hellings. 2014. Conjunctive context-free path queries. In *Proceedings of ICDT'14*. 119–130.
- [7] Jelle Hellings. 2015. Querying for Paths in Graphs using Context-Free Path Queries. *arXiv preprint arXiv:1502.02242* (2015).
- [8] Jochem Kuijpers, George Fletcher, Nikolay Yakovets, and Tobias Lindaaker. 2019. An Experimental Study of Context-Free Path Query Evaluation Methods. In *Proceedings of the 31st International Conference on Scientific and Statistical Database Management (SSDBM '19)*. ACM, New York, NY, USA, 121–132. <https://doi.org/10.1145/3335783.3335791>
- [9] Ciro M. Medeiros, Martin A. Musicante, and Umberto S. Costa. 2018. Efficient Evaluation of Context-free Path Queries for Graph Databases. In *Proceedings of the 33rd Annual ACM Symposium on Applied Computing (SAC '18)*. ACM, New York, NY, USA, 1230–1237. <https://doi.org/10.1145/3167132.3167265>
- [10] Nikita Mishin, Iaroslav Sokolov, Egor Spirin, Vladimir Kutuev, Egor Nemchinov, Sergey Gorbatyuk, and Semyon Grigorev. 2019. Evaluation of the Context-Free Path Querying Algorithm Based on Matrix Multiplication. In *Proceedings of the 2Nd Joint International Workshop on Graph Data Management Experiences & Systems (GRADES) and Network Data Analytics (NDA) (GRADES-NDA'19)*. ACM, New York, NY, USA, Article 12, 5 pages. <https://doi.org/10.1145/3327964.3328503>
- [11] Fred C. Santos, Umberto S. Costa, and Martin A. Musicante. 2018. A Bottom-Up Algorithm for Answering Context-Free Path Queries in Graph Databases. In *Web Engineering*, Tommi Mikkonen, Ralf Klamma, and Juan Hernández (Eds.). Springer International Publishing, Cham, 225–233.
- [12] Ekaterina Verbitskaia, Ilya Kirillov, Ilya Nozkin, and Semyon Grigorev. 2018. Parser Combinators for Context-free Path Querying. In *Proceedings of the 9th ACM SIGPLAN International Symposium on Scala (Scala 2018)*. ACM, New York, NY, USA, 13–23. <https://doi.org/10.1145/3241653.3241655>
- [13] Carl Yang, Aydin Buluc, and John D. Owens. 2019. GraphBLAST: A High-Performance Linear Algebra-based Graph Framework on the GPU. *arXiv:cs.DC/1908.01407*
- [14] Mihalis Yannakakis. 1990. Graph-theoretic methods in database theory. In *Proceedings of the ninth ACM SIGACT-SIGMOD-SIGART symposium on Principles of database systems*. ACM, 230–242.
- [15] X. Zhang, Z. Feng, X. Wang, G. Rao, and W. Wu. 2016. Context-free path queries on RDF graphs. In *International Semantic Web Conference*. Springer, 632–648.