

Title

Subtitile

First Author · Second Author

Received: date / Accepted: date

Abstract Insert your abstract here. Include keywords, PACS and mathematical subject classification numbers as needed.

Keywords First keyword · Second keyword · More

1 Introduction

2 The Motivating Example

In this section, we formulate the problem of context-free path query evaluation, using a small graph and the classical *same-generation query* [1], which cannot be expressed using regular expressions.

Let us have a graph database or any other object, which can be represented as a graph. The same-generation query can be used for discovering a vertex similarity, for example, gene similarity [2]. For graph databases, the same-generation query is aimed at the finding all the nodes at the same hierarchy level. The language, formed by the paths between such nodes, is not regular and corresponds to the language of matching parentheses. Hence, the query is formulated as a context-free grammar.

For example, let us have a small double-cyclic graph (see Figure 1). One of the cycles has three edges, labeled with a , and the other has two edges, labeled with b . Both cycles are connected via a shared node 0.

Supported by the Russian Science Foundation grant 18-11-00100.

F. Author
first address
Tel.: +123-45-678910
Fax: +123-45-678910
E-mail: fauthor@example.com

S. Author
second address




Fig. 1: An example graph.

For this graph, we have a same-generation query, formulated as a context-free grammar, which generates a context-free language $L = \{a^n b^n \mid n \geq 1\}$.

The result of context-free path query evaluation for this example is a set of node pairs (m, n) , such that there is a path from the node m to the node n , whose labeling forms a word from the language L . For example, the node pair $(0, 0)$ must be in this set, since there is a path from the node 0 to the node 0, whose labeling forms a string $w = aaaaaabbbbb = a^6 b^6 \in L$.

3 Preliminaries

Let Σ be a finite set of edge labels. Define an *edge-labeled directed graph* as a tuple $D = (V, E)$ with a set of nodes V and a directed edge relation $E \subseteq V \times \Sigma \times V$. For a path π in a graph D , we denote the unique word, obtained by concatenating the labels of the edges along the path π as $l(\pi)$. Also, we write $n\pi m$ to indicate, that the path π starts at the node $n \in V$ and ends at the node $m \in V$.

Following Hellings [3], we deviate from the usual definition of a context-free grammar in *Chomsky Normal Form* [4] by not including a special starting non-terminal, which will be specified in the path queries for the graph. Since every context-free grammar can be transformed into an equivalent one in Chomsky Normal Form and checking, that empty string belongs to the language is trivial, it is sufficient to consider only grammars of the following type. A *context-free grammar* is a triple $G = (N, \Sigma, P)$, where N is a finite set of non-terminals, Σ is a finite set of terminals, and P is a finite set of productions of the following forms:

- $A \rightarrow BC$, for $A, B, C \in N$,
- $A \rightarrow x$, for $A \in N$ and $x \in \Sigma$.

Note that we omit the rules of the form $A \rightarrow \varepsilon$, where ε denotes empty string. This does not restrict the applicability of our algorithm since only the empty paths $m\pi m$ correspond to empty string ε .

We use the conventional notation $A \xrightarrow{*} w$ to denote, that a string $w \in \Sigma^*$ can be derived from a non-terminal A by some sequence of production rule applications from P . The *language* of a grammar $G = (N, \Sigma, P)$ with respect to a start non-terminal $S \in N$ is defined by

$$L(G_S) = \{w \in \Sigma^* \mid S \xrightarrow{*} w\}.$$

For a given graph $D = (V, E)$ and a context-free grammar $G = (N, \Sigma, P)$, we define *context-free relations* $R_A \subseteq V \times V$ for every $A \in N$, such that

$$R_A = \{(n, m) \mid \exists n\pi m \ (l(\pi) \in L(G_A))\}.$$

We define a binary operation (\cdot) for arbitrary subsets N_1, N_2 of N with respect to a context-free grammar $G = (N, \Sigma, P)$ as

$$N_1 \cdot N_2 = \{A \mid \exists B \in N_1, \exists C \in N_2 \text{ such that } (A \rightarrow BC) \in P\}.$$

Using this binary operation as subset multiplication, and union as an addition, we can define a *matrix multiplication*, $a \times b = c$, where a and b are matrices of a suitable size, that have subsets of N as elements, as

$$c_{i,j} = \bigcup_{k=1}^n a_{i,k} \cdot b_{k,j}.$$

Also, we use the element-wise union operation on matrices a and b with the same size: $a \cup b = c$, where $c_{i,j} = a_{i,j} \cup b_{i,j}$.

According to Valiant [5], we define the *transitive closure* of a square matrix a as $a^+ = a^{(1)} \cup a^{(2)} \cup \dots$, where $a^{(1)} = a$ and

$$a^{(i)} = \bigcup_{j=1}^{i-1} a^{(j)} \times a^{(i-j)}, \quad i \geq 2.$$

Valiant proposed an algorithm for a context-free recognition for a linear input, which in graph terms corresponds to a directed chain of nodes. The algorithm enumerates the positions in the input string s from 0 to the length of s , constructs an upper-triangular matrix, and computes its transitive closure. In the context-free path querying input graphs can be arbitrary, which removes the upper-triangularity limitation. For this reason, we introduce another definition of transitive closure for arbitrary square matrix a as $a^{cf} = a^{(1)} \cup a^{(2)} \cup \dots$, where $a^{(1)} = a$ and

$$a^{(i)} = a^{(i-1)} \cup (a^{(i-1)} \times a^{(i-1)}), \quad i \geq 2.$$

These two transitive closure definitions are in fact equivalent (a formal proof can be found in Appendix ??). Furthermore, in this paper we use the transitive closure a^{cf} instead of a^+ , and the algorithm for computing a^{cf} also computes Valiant's transitive closure a^+ .

4 Related Work

Traditionally, query languages for graph databases use regular expressions to describe paths to find [6, ?, ?, ?, ?], but there are some other useful queries, which cannot be expressed by regular expressions. For example, there are classical *same-generation queries* [1], which can be used for finding all the nodes at the same level in some hierarchy, and are useful for discovering vertex similarity. The context-free path querying algorithms can be used to evaluate such types of queries since this queries can be represented by context-free grammars.

There are a number of solutions [3, ?, ?] for context-free path query evaluation w.r.t. relational query semantics, which make use of such parsing algorithms as CYK [7, ?] or Earley [8].

Hellings [3] presented an algorithm for context-free path query evaluation using relational query semantics. According to Hellings, for a given graph $D = (V, E)$ and a grammar $G = (N, \Sigma, P)$ the context-free path query evaluation w.r.t. relational query semantics reduces to a calculation of a set of context-free relations R_A . Thus, in this paper, we focus on the calculation of these context-free relations. Also, the algorithm in [3] was implemented by [9] in the context of RDF processing.

Other examples of path query semantics are *single-path* and *all-path query semantics* [10]. The all-path query semantics requires a finding of all possible paths from a node m to a node n whose labelings are derived from a non-terminal A . The single-path query semantics requires presenting only one such path. Hellings [10] presented some algorithms for context-free path query evaluation using single-path and all-path query semantics. If a context-free path query w.r.t. all-path query semantics is evaluated for cyclic graphs, then the query result can be an infinite set of paths. For this reason, in [10] annotated grammars were proposed as a way to represent the results.

In [11], an algorithm for a context-free path query evaluation w.r.t. all-path query semantics is proposed. This algorithm is based on the generalized top-down parsing algorithm (GLL) [12]. For the result representation, this solution uses derivation trees, which is more native for grammar-based analysis. The algorithms [11, ?] for context-free path query evaluation w.r.t. all-path query semantics can also be used for query evaluation using relational and single-path semantics.

Our work is inspired by Valiant [5], who proposed an algorithm for general context-free recognition in less than cubic time. This algorithm computes the same parsing table as CYK algorithm but does this by offloading the most intensive computations into calls to the Boolean matrix multiplication procedure. This approach not only provides an asymptotically more efficient algorithm but also allows us to effectively apply GPGPU computing techniques. Valiant's algorithm computes the transitive closure a^+ of a square upper-triangular matrix a . Valiant also has shown, that the matrix multiplication operation (\times) is essentially the same as $|N|^2$ Boolean matrix multiplications, where $|N|$ is the number of non-terminals in the given context-free grammar in Chomsky normal form.

Yannakakis [13] analyzed the reducibility of various path querying problems to the calculation of transitive closure. He formulated a problem of Valiant's technique generalization for the context-free path query evaluation w.r.t. relational query semantics. Also, he conjectured, that this technique cannot be generalized for arbitrary graphs, though it does for acyclic graphs.

Thus, reducing a context-free path query evaluation w.r.t. relational query semantics to a calculation of matrix transitive closure was an open problem until now.

5 GPU section

6 Context-free Path Querying by Transitive Closure Calculation

In this section, we show, how the context-free path query evaluation using relational query semantics can be reduced to the calculation of matrix transitive closure a^{cf} , prove the correctness of this reduction, introduce an algorithm for computing the transitive closure a^{cf} , and provide a step-by-step demonstration of this algorithm on a small example.

6.1 Reducing Context-Free Path Querying to the Calculation of Transitive Closure

In this section, we show, how the context-free relations R_A can be calculated by computing the transitive closure a^{cf} .

Let $G = (N, \Sigma, P)$ be a grammar and $D = (V, E)$ be a graph. We enumerate the nodes of the graph D from 0 to $(|V| - 1)$. We initialize the elements of the $|V| \times |V|$ matrix a with \emptyset . Further, for every i and j we set

$$a_{i,j} = \{A_k \mid ((i, x, j) \in E) \wedge ((A_k \rightarrow x) \in P)\}.$$

Finally, we compute the transitive closure

$$a^{cf} = a^{(1)} \cup a^{(2)} \cup \dots$$

where

$$a^{(i)} = a^{(i-1)} \cup (a^{(i-1)} \times a^{(i-1)}),$$

for $i \geq 2$ and $a^{(1)} = a$. For the transitive closure a^{cf} , the following statements hold.

Lemma 1 *Let $D = (V, E)$ be a graph, let $G = (N, \Sigma, P)$ be a grammar. Then for any i, j and for any non-terminal $A \in N$, $A \in a_{i,j}^{(k)}$ iff $(i, j) \in R_A$ and $i \pi j$, such that there is a derivation tree of the height $h \leq k$ for the string $l(\pi)$ and a context-free grammar $G_A = (N, \Sigma, P, A)$.*

Proof (Proof by Induction)

Base case: Show that the lemma holds for $k = 1$. For any i, j and for any non-terminal $A \in N$, $A \in a_{i,j}^{(1)}$ iff there is $i\pi j$ that consists of a unique edge e from the node i to the node j and $(A \rightarrow x) \in P$, where $x = l(\pi)$. Therefore $(i, j) \in R_A$ and there is a derivation tree of the height $h = 1$, shown on Figure 2, for the string x and a context-free grammar $G_A = (N, \Sigma, P, A)$. Thus, it has been shown that the lemma holds for $k = 1$.




Fig. 2: The derivation tree of the height $h = 1$ for the string $x = l(\pi)$.

Inductive step: Assume that the lemma holds for any $k \leq (p - 1)$ and show that it also holds for $k = p$, where $p \geq 2$. For any i, j and for any non-terminal $A \in N$,

$$A \in a_{i,j}^{(p)} \text{ iff } A \in a_{i,j}^{(p-1)} \text{ or } A \in (a^{(p-1)} \times a^{(p-1)})_{i,j},$$

since

$$a^{(p)} = a^{(p-1)} \cup (a^{(p-1)} \times a^{(p-1)}).$$

Let $A \in a_{i,j}^{(p-1)}$. By the inductive hypothesis, $A \in a_{i,j}^{(p-1)}$ iff $(i, j) \in R_A$ and there exists $i\pi j$, such that there is a derivation tree of the height $h \leq (p - 1)$ for the string $l(\pi)$ and a context-free grammar $G_A = (N, \Sigma, P, A)$. The statement of the lemma holds for $k = p$ since the height h of this tree is also less than or equal to p .

Let $A \in (a^{(p-1)} \times a^{(p-1)})_{i,j}$. By the definition of the binary operation (\cdot) on arbitrary subsets, $A \in (a^{(p-1)} \times a^{(p-1)})_{i,j}$ iff there are r , $B \in a_{i,r}^{(p-1)}$ and $C \in a_{r,j}^{(p-1)}$, such that $(A \rightarrow BC) \in P$. Hence, by the inductive hypothesis, there are $i\pi_1 r$ and $r\pi_2 j$, such that $(i, r) \in R_B$ and $(r, j) \in R_C$, and there are the derivation trees T_B and T_C of heights $h_1 \leq (p - 1)$ and $h_2 \leq (p - 1)$ for the strings $w_1 = l(\pi_1)$, $w_2 = l(\pi_2)$ and the context-free grammars G_B , G_C respectively. Thus, the concatenation of paths π_1 and π_2 is $i\pi j$, where $(i, j) \in R_A$ and there is a derivation tree of the height $h = 1 + \max(h_1, h_2)$, shown on Figure 3, for the string $w = l(\pi)$ and a context-free grammar G_A .

The statement of the lemma holds for $k = p$ since the height $h = 1 + \max(h_1, h_2) \leq p$. This completes the proof of the lemma.




Fig. 3: The derivation tree of the height $h = 1 + \max(h_1, h_2)$ for the string $w = l(\pi)$, where T_B and T_C are the derivation trees for strings w_1 and w_2 respectively.

Theorem 1 Let $D = (V, E)$ be a graph and let $G = (N, \Sigma, P)$ be a grammar. Then for any i, j and for any non-terminal $A \in N$, $A \in a_{i,j}^{cf}$ iff $(i, j) \in R_A$.

Proof Since the matrix $a^{cf} = a^{(1)} \cup a^{(2)} \cup \dots$, for any i, j and for any non-terminal $A \in N$, $A \in a_{i,j}^{cf}$ iff there is $k \geq 1$, such that $A \in a_{i,j}^{(k)}$. By the lemma 1, $A \in a_{i,j}^{(k)}$ iff $(i, j) \in R_A$ and there is $i\pi j$, such that there is a derivation tree of the height $h \leq k$ for the string $l(\pi)$ and a context-free grammar $G_A = (N, \Sigma, P, A)$. This completes the proof of the theorem.

We can, therefore, determine whether $(i, j) \in R_A$ by asking whether $A \in a_{i,j}^{cf}$. Thus, we show how the context-free relations R_A can be calculated by computing the transitive closure a^{cf} of the matrix a .

6.2 The algorithm

In this section, we introduce an algorithm for calculating the transitive closure a^{cf} which was discussed in section 6.1.

Let $D = (V, E)$ be the input graph and $G = (N, \Sigma, P)$ be the input grammar.

Algorithm 1 Context-free recognizer for graphs

```

1: function CONTEXTFREEPATHQUERYING( $D, G$ )
2:    $n \leftarrow$  the number of nodes in  $D$ 
3:    $E \leftarrow$  the directed edge-relation from  $D$ 
4:    $P \leftarrow$  the set of production rules in  $G$ 
5:    $T \leftarrow$  the matrix  $n \times n$  in which each element is  $\emptyset$ 
6:   for all  $(i, x, j) \in E$  do                                      $\triangleright$  Matrix initialization
7:      $T_{i,j} \leftarrow T_{i,j} \cup \{A \mid (A \rightarrow x) \in P\}$ 
8:   while matrix  $T$  is changing do
9:      $T \leftarrow T \cup (T \times T)$                                     $\triangleright$  Transitive closure  $T^{cf}$  calculation
10:  return  $T$ 

```

Note, the matrix initialization in lines **6-7** of the Algorithm 1 can handle arbitrary graph D . For example, if a graph D contains multiple edges (i, x_1, j)

and (i, x_2, j) then both the elements of the set $\{A \mid (A \rightarrow x_1) \in P\}$ and the elements of the set $\{A \mid (A \rightarrow x_2) \in P\}$ will be added to $T_{i,j}$.

We need to show that the Algorithm 1 terminates in a finite number of steps. Since each element of the matrix T contains no more than $|N|$ non-terminals, the total number of non-terminals in the matrix T does not exceed $|V|^2|N|$. Therefore, the following theorem holds.

Theorem 2 *Let $D = (V, E)$ be a graph and let $G = (N, \Sigma, P)$ be a grammar. The Algorithm 1 terminates in a finite number of steps.*

Proof It is sufficient to show, that the operation in the line 9 of the Algorithm 1 changes the matrix T only finite number of times. Since this operation can only add non-terminals to some elements of the matrix T , but not remove them, it can change the matrix T no more than $|V|^2|N|$ times.

Denote the number of elementary operations executed by the algorithm of multiplying two $n \times n$ Boolean matrices as $BMM(n)$. According to Valiant, the matrix multiplication operation in the line 9 of the Algorithm 1 can be calculated in $O(|N|^2 BMM(|V|))$. Denote the number of elementary operations, executed by the matrix union operation of two $n \times n$ Boolean matrices as $BMU(n)$. Similarly, it can be shown that the matrix union operation in the line 9 of the Algorithm 1 can be calculated in $O(|N|^2 BMU(n))$. Since the line 9 of the Algorithm 1 is executed no more than $|V|^2|N|$ times, the following theorem holds.

Theorem 3 *Let $D = (V, E)$ be a graph and let $G = (N, \Sigma, P)$ be a grammar. The Algorithm 1 calculates the transitive closure T^{cf} in $O(|V|^2|N|^3(BMM(|V|) + BMU(|V|)))$.*

We also provide the worst-case example, for which the time complexity in terms of the graph size provided by Theorem 3 cannot be improved. This example is based on the context-free grammar $G = (N, \Sigma, P)$ where:

- the set of non-terminals $N = \{S\}$;
- the set of terminals $\Sigma = \{a, b\}$;
- the set of production rules P is presented on Figure 4.

$$\begin{array}{l} 0 : S \rightarrow a S b \\ 1 : S \rightarrow a b \end{array}$$

Fig. 4: Production rules for the worst-case example.

Let the size $|N|$ of the grammar G be a constant. The worst-case time complexity is reached by running this query on the double-cyclic graph where:

- one of the cycles having $u = 2^k + 1$ edges labeled with a ;
- another cycle having $v = 2^k$ edges labeled with b ;

- the two cycles are connected via a shared node m .

A small example of such graph with $k = 1$, $u = 3$, $v = 2$, and $m = 0$ is presented on Figure 1.

The shortest path π from the node m to the node m , whose labeling forms a string from the language $L(G_S) = \{a^n b^n | n \geq 1\}$, has a length $l = 2 * u * v$, since $u = 2^k + 1$ and $v = 2^k$ are coprime, and string s , formed by this path, consists of $u * v$ labels a and $u * v$ labels b . The string $s = l(\pi)$ has a derivation tree according to a context-free grammar G_S of the minimal height $h = 2 * u * v$ among all the paths from the node m to the node m in this double-cyclic graph. Therefore, if we run the worst-case example query on this graph, then the operation in the line **9** of the Algorithm 1 changes the matrix T at least $h = 2 * u * v$ times. Hence, the Algorithm 1 computes this query in $O(|V|^2(BMM(|V|) + BMU(|V|)))$, since $|V| = (u + v - 1) = 2 * v$ and $h = 2 * u * v > 2 * v * v = |V|^2/4 = O(|V|^2)$.

6.3 An Example

In this section, we provide a step-by-step demonstration of the proposed algorithm. For this, we consider the example with the worst-case time complexity.

The **example query** is based on the context-free grammar $G = (N, \Sigma, P)$ of the worst-case example query which was discussed in section 6.2. The set of production rules for this grammar is shown on Figure 4.

Since the proposed algorithm processes only grammars in Chomsky normal form, we first transform the grammar G into an equivalent grammar $G' = (N', \Sigma', P')$ in normal form, where:

- the set of non-terminals $N' = \{S, S_1, A, B\}$;
- the set of terminals $\Sigma' = \{a, b\}$;
- the set of production rules P' is presented on Figure 5.

$$\begin{aligned} 0 : S &\rightarrow A B \\ 1 : S &\rightarrow A S_1 \\ 2 : S_1 &\rightarrow S B \\ 3 : A &\rightarrow a \\ 4 : B &\rightarrow b \end{aligned}$$

Fig. 5: Production rules for the example query grammar in normal form.

We run the query on a graph, presented on Figure 1. We provide a step-by-step demonstration of the work with the given graph D and grammar G' of the Algorithm 1. After the matrix initialization in lines **6-7** of the Algorithm 1, we have a matrix T_0 , presented on Figure 6.

Let T_i be the matrix T , obtained after executing the loop in the lines **8-9** of the Algorithm 1 i times. The calculation of the matrix T_1 is shown on Figure 7.

$$T_0 = \begin{pmatrix} \emptyset & \{A\} & \emptyset & \{B\} \\ \emptyset & \emptyset & \{A\} & \emptyset \\ \{A\} & \emptyset & \emptyset & \emptyset \\ \{B\} & \emptyset & \emptyset & \emptyset \end{pmatrix}$$

Fig. 6: The initial matrix for the example query.

$$T_0 \times T_0 = \begin{pmatrix} \emptyset & \emptyset & \emptyset & \emptyset \\ \emptyset & \emptyset & \emptyset & \emptyset \\ \emptyset & \emptyset & \emptyset & \{S\} \\ \emptyset & \emptyset & \emptyset & \emptyset \end{pmatrix}$$

$$T_1 = T_0 \cup (T_0 \times T_0) = \begin{pmatrix} \emptyset & \{A\} & \emptyset & \{B\} \\ \emptyset & \emptyset & \{A\} & \emptyset \\ \{A\} & \emptyset & \emptyset & \{S\} \\ \{B\} & \emptyset & \emptyset & \emptyset \end{pmatrix}$$

Fig. 7: The first iteration of computing the transitive closure for the example query.

When the algorithm at some iteration finds new paths in the graph D , then it adds corresponding nonterminals to the matrix T . For example, after the first loop iteration, non-terminal S is added to the matrix T . This non-terminal is added to the element with a row index $i = 2$ and a column index $j = 3$. This means, that there is $i\pi j$ (a path π from the node 2 to the node 3), such that $S \xrightarrow{*} l(\pi)$. For example, such a path consists of two edges with labels a and b , and thus $S \xrightarrow{*} a b$.

The calculation of the transitive closure is completed after k iterations, when a fixpoint is reached: $T_{k-1} = T_k$. For the example query, $k = 13$ since $T_{13} = T_{12}$. The remaining iterations of computing the transitive closure are presented on Figure 8 (new matrix elements on each iteration are shown in bold).

Thus, the result of the Algorithm 1 for the example query is the matrix $T_{13} = T_{12}$. Now, after constructing the transitive closure, we can construct the context-free relations R_A . These relations for each non-terminal of the grammar G' are presented on Figure 9.

In the context-free relation R_S , we have all node pairs corresponding to paths, whose labeling is in the language $L(G_S) = \{a^n b^n | n \geq 1\}$. This conclusion is based on the fact, that the grammar G'_S is equivalent to the grammar G_S and $L(G_S) = L(G'_S)$.

$$\begin{aligned}
T_2 &= \begin{pmatrix} \emptyset & \{A\} & \emptyset & \{B\} \\ \emptyset & \emptyset & \{A\} & \emptyset \\ \{A, \mathbf{S_1}\} & \emptyset & \emptyset & \{S\} \\ \{B\} & \emptyset & \emptyset & \emptyset \end{pmatrix} & T_3 &= \begin{pmatrix} \emptyset & \{A\} & \emptyset & \{B\} \\ \{\mathbf{S_1}\} & \emptyset & \{A\} & \emptyset \\ \{A, S_1\} & \emptyset & \emptyset & \{S\} \\ \{B\} & \emptyset & \emptyset & \emptyset \end{pmatrix} & T_4 &= \begin{pmatrix} \emptyset & \{A\} & \emptyset & \{B\} \\ \{S\} & \emptyset & \{A\} & \{\mathbf{S_1}\} \\ \{A, S_1\} & \emptyset & \emptyset & \{S\} \\ \{B\} & \emptyset & \emptyset & \emptyset \end{pmatrix} \\
T_5 &= \begin{pmatrix} \emptyset & \{A\} & \emptyset & \{B, \mathbf{S_1}\} \\ \{S\} & \emptyset & \{A\} & \{S_1\} \\ \{A, S_1\} & \emptyset & \emptyset & \{S\} \\ \{B\} & \emptyset & \emptyset & \emptyset \end{pmatrix} & T_6 &= \begin{pmatrix} \{\mathbf{S_1}\} & \{A\} & \emptyset & \{B, S\} \\ \{S\} & \emptyset & \{A\} & \{S_1\} \\ \{A, S_1\} & \emptyset & \emptyset & \{S\} \\ \{B\} & \emptyset & \emptyset & \emptyset \end{pmatrix} & T_7 &= \begin{pmatrix} \{S_1\} & \{A\} & \emptyset & \{B, S\} \\ \{S\} & \emptyset & \{A\} & \{S_1\} \\ \{A, S_1, \mathbf{S_1}\} & \emptyset & \emptyset & \{S\} \\ \{B\} & \emptyset & \emptyset & \emptyset \end{pmatrix} \\
T_8 &= \begin{pmatrix} \{S_1\} & \{A\} & \emptyset & \{B, S\} \\ \{S\} & \emptyset & \{A\} & \{S_1\} \\ \{A, S_1, S\} & \emptyset & \emptyset & \{S, \mathbf{S_1}\} \\ \{B\} & \emptyset & \emptyset & \emptyset \end{pmatrix} & T_9 &= \begin{pmatrix} \{S_1\} & \{A\} & \emptyset & \{B, S\} \\ \{S\} & \emptyset & \{A\} & \{S_1, \mathbf{S_1}\} \\ \{A, S_1, S\} & \emptyset & \emptyset & \{S, S_1\} \\ \{B\} & \emptyset & \emptyset & \emptyset \end{pmatrix} & T_{10} &= \begin{pmatrix} \{S_1\} & \{A\} & \emptyset & \{B, S\} \\ \{S, \mathbf{S_1}\} & \emptyset & \{A\} & \{S_1, S\} \\ \{A, S_1, S\} & \emptyset & \emptyset & \{S, S_1\} \\ \{B\} & \emptyset & \emptyset & \emptyset \end{pmatrix} \\
T_{11} &= \begin{pmatrix} \{S_1, \mathbf{S_1}\} & \{A\} & \emptyset & \{B, S\} \\ \{S, S_1\} & \emptyset & \{A\} & \{S_1, S\} \\ \{A, S_1, S\} & \emptyset & \emptyset & \{S, S_1\} \\ \{B\} & \emptyset & \emptyset & \emptyset \end{pmatrix} & T_{12} &= \begin{pmatrix} \{S_1, S\} & \{A\} & \emptyset & \{B, S, \mathbf{S_1}\} \\ \{S, S_1\} & \emptyset & \{A\} & \{S_1, S\} \\ \{A, S_1, S\} & \emptyset & \emptyset & \{S, S_1\} \\ \{B\} & \emptyset & \emptyset & \emptyset \end{pmatrix} & T_{13} &= \begin{pmatrix} \{S_1, S\} & \{A\} & \emptyset & \{B, S, S_1\} \\ \{S, S_1\} & \emptyset & \{A\} & \{S_1, S\} \\ \{A, S_1, S\} & \emptyset & \emptyset & \{S, S_1\} \\ \{B\} & \emptyset & \emptyset & \emptyset \end{pmatrix}
\end{aligned}$$

Fig. 8: Remaining states of the matrix T .

$$\begin{aligned}
R_S &= \{(0, 0), (0, 3), (1, 0), (1, 3), (2, 0), (2, 3)\}, \\
R_{S_1} &= \{(0, 0), (0, 3), (1, 0), (1, 3), (2, 0), (2, 3)\}, \\
R_A &= \{(0, 1), (1, 2), (2, 0)\}, \\
R_B &= \{(0, 3), (3, 0)\}.
\end{aligned}$$

Fig. 9: Context-free relations for the example query.

7 Evaluation

8 Conclusion

References

1. S. Abiteboul, R. Hull, V. Vianu, *Foundations of databases: the logical level* (Addison-Wesley Longman Publishing Co., Inc., 1995)
2. P. Sevon, L. Eronen, *Journal of Integrative Bioinformatics* **5**(2), 100 (2008)
3. J. Hellings, in *Proceedings of ICDT'14* (2014), pp. 119–130
4. N. Chomsky, *Information and control* **2**(2), 137 (1959)
5. L.G. Valiant, *Journal of computer and system sciences* **10**(2), 308 (1975)
6. J.L. Reutter, M. Romero, M.Y. Vardi, *Theory of Computing Systems* **61**(1), 31 (2017)
7. T. Kasami, An efficient recognition and syntaxanalysis algorithm for context-free languages. Tech. rep., DTIC Document (1965)
8. D. Grune, C.J.H. Jacobs, *Parsing Techniques (Monographs in Computer Science)* (Springer-Verlag New York, Inc., Secaucus, NJ, USA, 2006)
9. X. Zhang, Z. Feng, X. Wang, G. Rao, W. Wu, in *International Semantic Web Conference* (Springer, 2016), pp. 632–648
10. J. Hellings, arXiv preprint arXiv:1502.02242 (2015)
11. S. Grigorev, A. Ragoza, arXiv preprint arXiv:1612.08872 (2016)
12. E. Scott, A. Johnstone, *Electronic Notes in Theoretical Computer Science* **253**(7), 177 (2010)

13. M. Yannakakis, in *Proceedings of the ninth ACM SIGACT-SIGMOD-SIGART symposium on Principles of database systems* (ACM, 1990), pp. 230–242