

Matrix-based path querying algorithms on the GPU

Rustam Azimov

Keywords

CFPQ, context-free path querying, matrix operations, GPGPU, sparse matrices, formal language theory, context-free grammars.

1 Introduction

Graphs are used as a data structure to represent large volumes of information in a compact and convenient for analysis form in many areas: graph databases [1], bioinformatics [2], static code analysis [3, 4], etc. In these areas, it is necessary to evaluate queries for large graphs in order to determine the dependencies between the nodes.

The most popular are queries that use regular or context-free grammars for constraints on the path. The use of context-free instead of regular grammars in path querying allows us to formulate more complex queries to the graph and solve a wider class of problems.

There is a number of algorithms for context-free path query evaluation [5, 6, 7, 8]. However, these algorithms demonstrate a poor performance when applied to large graphs.

One of the most popular techniques used to increase performance when working with large data is the use of a GPU to perform computation. We proposed the only known approach in this area that makes it possible to use the GPUs effectively [9]. This is the matrix approach, in which the adjacency matrix of the input graph is built, the elements of which are sets of non-terminals of the input grammar. Next, the transitive closure of the constructed matrix is calculated using the derivation rules of the input grammar. In the process of computing, the operations of multiplication and addition of Boolean matrices are actively used and can be computed very efficiently using the GPUs. For example, such libraries as cuBLAS or cuSPARSE provide high-performance implementations of necessary matrix operations (GEMM, GEAM).

The project is devoted to the study of our new matrix-based path querying approach and creating a high-performance implementations for analyzing large graphs. The active use of matrix operations in our approach makes it possible to efficiently apply a wide class of matrix optimizations and computing techniques (GPGPU, parallel processing, sparse matrix representation, distributed-memory computation, etc.).

2 Motivation

There are other projects which are interested in a high-performance context-free path querying implementations.

For example, there is the GraphBLAS that provides standardized building blocks of graph algorithms, optimizes and simplifies computation of many different graph queries. We can provide a high-performance context-free path querying implementations as part of GraphBLAS. In these implementations we can use sparse matrix representation and with the help of Nvidia libraries we can apply such optimizations as GPGPU and distributed-memory computation.

Also, there is the industry’s most widely adopted graph database query language Cypher. The context-free path querying can be applied very efficiently to the graph databases. Therefore, Cypher is interested in supporting context-free path querying and having a high-performance implementations for these queries.

In addition, there is the RedisGraph which is the fastest graph database with the smallest storage footprint. It analyze and connect millions of distributed data points in real time and deliver contextual content to customers. RedisGraph uses both GraphBLAS and Cypher Query. Therefore, we also can contribute a high-performance implementations of the expressive graph database queries to the RedisGraph.

References

- [1] A. Mendelzon and P. Wood. Finding regular simple paths in graph databases. *SIAM J. Computing*, 24(6):1235–1258, 1995.
- [2] James WJ Anderson, Ádám Novák, Zsuzsanna Sükösd, Michael Golden, Preeti Arunapuram, Ingolfur Edvardsson, and Jotun Hein. Quantifying variances in comparative rna secondary structure prediction. *BMC bioinformatics*, 14(1):149, 2013.
- [3] John Kodumal and Alex Aiken. The set constraint/cfl reachability connection in practice. *ACM Sigplan Notices*, 39(6):207–218, 2004.
- [4] Qirun Zhang, Michael R Lyu, Hao Yuan, and Zhendong Su. Fast algorithms for dyck-cfl-reachability with applications to alias analysis. In *ACM SIGPLAN Notices*, volume 48, pages 435–446. ACM, 2013.
- [5] Semyon Grigorev and Anastasiya Ragozina. Context-free path querying with structural representation of result. *arXiv preprint arXiv:1612.08872*, 2016.
- [6] Jelle Hellings. Conjunctive context-free path queries. In *Proceedings of ICDT’14*, pages 119–130, 2014.
- [7] X. Zhang, Z. Feng, X. Wang, G. Rao, and W. Wu. Context-free path queries on rdf graphs. In *International Semantic Web Conference*, pages 632–648. Springer, 2016.

- [8] Petteri Sevon and Lauri Eronen. Subgraph queries by context-free grammars. *Journal of Integrative Bioinformatics*, 5(2):100, 2008.
- [9] Rustam Azimov and Semyon Grigorev. Context-free path querying by matrix multiplication. In *Proceedings of the 1st ACM SIGMOD Joint International Workshop on Graph Data Management Experiences & Systems (GRADES) and Network Data Analytics (NDA)*, page 5. ACM, 2018.