Parser-Combinators for Context-Free Path Querying

Sophia Smolina
Electrotechnical University
St. Petersburg, Russia
sofysmol@gmail.com

Ekaterina Verbitskaia
Saint Petersburg State University
St. Petersburg, Russia
kajigor@gmail.com

Ilya Kirillov
Saint Petersburg State University
St. Petersburg, Russia
kirillov.ilija@gmail.com

Ilya Nozkin
Saint Petersburg State University
St. Petersburg, Russia
nozhkin.ii@gmail.com

Semyon Grigorev
Saint Petersburg State University
St. Petersburg, Russia
s.v.grigoriev@spbu.ru

ABSTRACT

A transparent integration of a domain-specific language for specification of context-free path queries (CFPQs) into a general-purpose programming language as well as static checking of errors in queries may greatly simplify the development of applications utilizing CF-PQs. Such techniques as LINQ and ORM can be used for the integration, but they have issues with flexibility: query decomposition and reusing of subqueries are a challenge. Adaptation of parser combinators technique for paths querying may solve these problems. Conventional parser combinators process linear input and only the Trails library is known to apply this technique for path querying. Trails suffers the common parser combinators issue: it does not support left-recursive grammars and also experiences problems in cycles handling. We demonstrate that it is possible to create general parser combinators for CFPQ which support arbitrary context-free grammars and arbitrary input graphs. We implement a library of such parser combinators and show that it is applicable for realistic tasks.

CCS CONCEPTS

• **Information systems → Graph-based database models**; **Query languages for non-relational engines**; • **Software and its engineering** → *Functional languages*; • **Theory of computation** → *Grammars and context-free languages*;

KEYWORDS

Graph Databases, Language-Constrained Path Problem, Context-Free Path Querying, Contex-Free Language Reachability, Parser Combinators, Generalized LL, GLL, Neo4J, Scala

1 INTRODUCTION

Graph querying is finding all paths in the graph which satisfy some constraints. If the constraints are specified with some formal language formalism, i.e. a grammar, it is called a language-constrained path query. For example, a grammar $S \rightarrow a\,S\,b \mid a\,b$ can be regarded as a query for the paths of the form $a^n b^n$, where $n \geq 1$. Querying the graph in the figure 1 returns the set of paths, each of which starts and ends in the vertex 3 and goes around the cycles in the graph the appropriate number of times: 3, 1, 2, 3, 1, 2, 3, 4, 3, 4, 3, 4, 3 and so on.


Figure 1: An example of input graph

Most existing graph traversing/querying languages, including SPARQL [21], Cypher [1], and Gremlin [26] support only regular languages as queries. For some applications regular languages are not expressive enough. Context-free path queries (CFPQ), on which we focus this paper, employ context-free languages for constraints specification. CFPQs are used in bioinformatics [30], static code analysis [2, 20, 23, 32, 36], and RDF processing [35]. Although there is a lot of problem-specific solutions and theoretical research on CFPQs [1, 6–8, 17, 25, 30, 33], cfSPARQL [35] is the single known graph query language to support CF constraints. Generic solution for the integration of CFPQs into general-purpose languages is not discussed enough.

When developing a data-centric application, one wants to use a general-purpose programming language and also to have a transparent and native access to data sources. One way to achieve this goal is to use string-embedded DSLs. In this approach, a query is written as a string, then passed on to a dedicated driver which executes it and returns a possibly untyped result. Despite the simplicity, string-embedded DSLs have serious drawbacks. First of all, they require the developer to learn the language itself, its features, runtime

---

[1] Cypher language web page: https://neo4j.com/developer/cypher-query-language/. Access date: 16.01.2018

and how the integration between the languages is implemented. DSLs are also a source of possible errors and vulnerabilities, static detection of which is a serious challenge [3]. Such techiques as the Object Relationship Mapping (ORM) or Language Integrated Query (LINQ) [4, 13, 16] partly solve these problems, but they still have issues with flexibility: both query decomposition and reusing of subqueries are a struggle. In this paper, we propose a transparent and natural integration of CFPQs into a general-purpose programming language.

Context-free path queries are known in various domains under different names. The *context-free language reachability framework* or *IFDS framework* is how it is called in the area of static code analysis. In [23, 24] Thomas Reps shows that the wide range of static code anlysis problems can be formulated in terms of CFL-reachability in the graph. This framework is used for such problems as the taint analysis [9], the alias analysis [31, 32, 36], the label flow analysis [20], and the fix locations problem [5]. What we propose in the paper can be used as a core of such framework since it provides both problem and domain independent mechanism for CFPQs evaluation.

We view parser combinators as the best way to integrate context-free language specifications into a general-purpose programming language. Parser combinators not only provide a transparent integration, but also compile-time checks of correctness and high-level techniques for generalization. An idea to use combinators for graph traversing has already been proposed in [12]. Unfortunately, the solution presented processes cycles in the input graph only approximately and is unable to handle left-recursive combinators, which is the most common issue of the approach. Authors pointed out that the idea described is similar to the classical parser combinators, but the language class supported or restrictions are not discussed.

Parser combinators are known to handle only a subset of context-free grammars: left recursion and ambiguity of the grammars are problematic. In [10], authors demonstrate a set of parser combinators which handles arbitrary context-free grammars by using ideas of the Generalized LL [27] algorithm (GLL). Meerkat [2] parser combinators library implements [10] and provides the parsing result in a compact form as the Shared Packed Parse Forest [22] (SPPF). SPPF is a suitable finite structural representation of a CFPQ result, even when the set of paths is infinite [6]. All the paths can be extracted from the SPPF—as the corresponding derivation trees—and further analysis can be done. It is also possible to run some further processing over the SPPF itself—without explicit paths extraction.

In this paper, we compose these ideas and present a set of parser-combinators for context-free path querying which handle arbitrary context-free grammars and provide a structural representation of the result. We make the following contributions in the paper.

(1) We show that it is possible to create a set of parser combinators for context-free path querying which work on both arbitrary context-free grammars and arbitrary graphs and provide a finite structural representation of the query result.

(2) We implement the parser combinators library in Scala. This library provides an integration to Neo4j [3] graph database.

The source code is available on GitHub: https://github.com/YaccConstructor/Meerkat.

(3) We perform an evaluation on realistic data and compare the performance of our library with another GLL-based CFPQ tool and with the Trails library. We conclude that our solution is expressive and performant enough to be applied to the real-world problems.

This paper is orgatized as follows. We introduce a fromal definition of the CFPQ problem in the section 2 and we provide a basic description of the Meerkat library and SPPF data structure in the section 3. We describe our solution in the section 4. In the section 5 we present and discuss a set of classical queries (the same generation query, the queries to a movie dataset [4]) formulated in terms of our library. Evaluation of the library is described in the section 6. Finally, in the section 7 we conclude and discuss possible directions of further research.

2 CONTEX-FREE PATH QUERYING PROBLEM

In this section we formally describe the context-free path querying problem (or context-free reachability problem).

First, we introduce the necessary definitions.

- Context-free grammar is a quadruple $G = (N, \Sigma, P, S)$, where N is a set of nonterminal symbols, Σ is a set of terminal symbols, $S \in N$ is a start nonterminal, and P is a set of productions.
- $\mathcal{L}(G)$ denotes a language specified by the grammar G, and is a set of terminal strings derived from the start nonterminal of G: $L(G) = \{\omega | S \Rightarrow^*_G \omega\}$.
- Directed graph is a triple $M = (V, E, L)$, where V is a set of vertices, $L \subseteq \Sigma$ is a set of labels, and a set of edges $E \subseteq V \times L \times V$. We assume that there are no parallel edges with equal labels: for every $e_1 = (v_1, l_1, v_2) \in E, e_2 = (u_1, l_2, u_2) \in E$ if $v_1 = u_1$ and $v_2 = u_2$ then $l_1 \neq l_2$.
- $tag : E \rightarrow L$ is a helper function which returns a tag of a given edge.

$$tag(e = (v_1, l, v_2), e \in E) = l$$

- $\oplus : L^+ \times L^+ \rightarrow L^+$ denotes a tag concatenation operation.
- Ω is a helper function which constructs a string produced by the given path. For every p path in M

$$\Omega(p = e_0, e_1, \ldots, e_{n-1}) = tag(e_0) \oplus \cdots \oplus tag(e_{n-1}).$$

We define the context-free language constrained path querying as, given a query in the form of a grammar G, to construct the set of the paths

$$Q(M, G) = \{p | p \text{ is path in } M, \Omega(p) \in \mathcal{L}(G)\}.$$

The CFL reachability problem is pretty similar and is formulated as follows:

$$Q(M, G) = \{(v_0, v_n) \mid \exists p \text{ is path in } M, p = v_0 \xrightarrow{l_0} \cdots \xrightarrow{l_{n-1}} v_n,$$
$$\Omega(p) \in \mathcal{L}(G)\}$$

.

Note that $Q(M, G)$ can be an infinite set, hence it cannot be represented explicitly. We show how to construct a compact data

structure which stores all the elements of $Q(M, G)$ in a finite space; every path can be extracted from this representation.

3 GENERALIZED PARSER COMBINATORS

Combinators techniques are shown to be applicable for graph traversing [12], but it still suffers the common issue with left-recursive definitions. A general parser combinators library Meerkat [10], implemented in the Scala programming language, removes this restrction by using memoization, continuation-passing style, and the ideas of Juhnson [11]. It supports the arbitrary (left-recursive and ambiguous) context-free specifications, but it also supports the specification of action code, and provide a syn macro for custom handling of the recursive nonterminal descriptions. Meerkat constructs the compact representation of the parse forest in the form of SPPF, which can be used for CFPQs results representation [6]. The worst case time and space complexity of the solution is cubic.

A Meerkat specification of the language $\{a^n b^n \mid n \geq 1\}$ is presented in fig. 2.

```
val S = syn("a" ~ S.? ~ "b")
```

Figure 2: Meercat specification of $\{a^n b^n \mid n \geq 1\}$

[6] showed that Generalized LL parsing algorithm [27] can be generalized to effectively process CFPQs and the query result can be finitely represented. As the Meerkat library is closely related to the Generalized LL algorithm and since GLL can be generalized for context-free path querying, it is also possible to adapt the Meerkat library for graph querying. It can be done by providing a function for retrieving the symbols which follow the specified position and utilizing it in the basic set of combinators. Details described below.

3.1 SPPF

Parsing of a string with respect to an ambiguous grammar can result in several derivation trees for a single string. The set of derivation trees is named a *derivation forest*. To store a derivation forest efficiently, the generalized parsing algorithms utilize a *Shared Packed Parse Forest* proposed by Joan Rekers [22]. The most efficient compact representation of derivation forests is a Binarized Shared Packed Parse Forest (we will abbreviate it to SPPF) [29]. GLL algorithm which utilizes this structure achieves the worst-case cubic space complexity [28].

Binarized SPPF is a directed graph, each node of which has one of the four types described below. If i is a start position of a substring and j—its end position, then (i, j) is called an *extension* of a node.

- **Terminal node** labeled (T, i, j).
- **Nonterminal node** labeled (N, i, j). This node denotes that there is at least one derivation $N \Rightarrow_G^* \omega[i \ldots j - 1]$—a substring of the input from i-th to j-th position. Every derivation tree for the given substring and nonterminal can be extracted by left-to-right top-down traversal of SPPF started from the respective node.
- **Intermediate node**: a special kind of node used for the binarization of the SPPF. These nodes are labeled with (t, i, j), where t is a grammar slot.

- **Packed node** labeled $(N \rightarrow \alpha, k)$, where k is a position in input of the right end of leftmost subtree of this node. A subgraph with the "root" in such node is one derivation from the nonterminal N.

An example of SPPF is presented in figure 4. We removed redundant intermediate and packed nodes for simplicity and to decrease the size of the figure.

SPPF can finitely represent a possibly infinite set of paths in the context of the language-constrained graph querying [6]. Since the SPPF stores derivation trees for all paths, it can be useful for the postprocessing and further understanding of the query results. In static code analysis, for example, it is possible to map paths back onto the source code thus providing a human-readable result.

4 PARSER COMBINATORS FOR PATH QUERYING

Parser combinators is a way to specify both a language syntax and a parser for it in terms of higher-order functions. Parser in this framework is a function which consumes a prefix of an input and returns either a parsing result or an error if the input is erroneous. Parser combinators compose parsers to form more complex parsers. A parser combinators library usually provides a set of a basic parser combinators, such as a combinator of the sequential application of the choice, but there can also be user-defined combinators. Most parser combinators libraries, including the Meerkat library, can only process the linear input—strings or some kind of streams. We modify the Meerkat library to work on the graph input.

The following ideas are at the core of the modification.

- The intersection of a context-free and a regular language is context-free. There are several constructive proofs of this fact. The proposed solution is a yet another constructive proof with the SPPF as a user-friendly representation of the context-free grammar for the intersection.
- Linear input can be regarded as a linear directed graph with symbols of the input labeling the edges.
- A conventional parser moves a pointer in the input from the position i to the position $i + 1$ and creates a new state when token between i-th and $i + 1$-th positions matches what is required in the grammar. In case of graph processing, there are possibly multiple ways to move from the current vertex i and it is possible to produce multiple new states. Generalized parsing is designed to optimally handle the production of multiple new states thus it is suitable to handle graph processing.
- Matching a token in the input can be viewed as a predicate, for example $p_c(x) = x == c$. We can generalize this observation allowing matching of an edge label of an arbitrary type with a predicate of some sort.
- If vertices of the graph contain any data of interest, we can treat them in the similar fashion as the edges. We can also convert the input graph transforming vertices into edges and then querying the transformed graph.

Querying process in our library is inherited from generalized parsing and is done in two steps. The first step is "parsing": the construction of the SPPF which contains all derivation trees for the paths satisfying syntactic constraints. The second step is semantic

actions application which retrieves the necessary additional data about the paths from the SPPF.

4.1 The Set of Combinators

We will demonstrate the set of combinators by the example: the input graph, which represent a map, is presented in fig. 3. There are some cities connected by one-way roads represented by the edges labeled *road_to*. Each city is labeled by its name and a country it belongs to.

**Figure 3: Example graph. Vertex labels are in the form "city-name (country-name)"**

Two basic building blocks of queries are the combinators for dealing with edges and vertices.

- `V[L](predicate: L =>Boolean)` the conbinator for processing vertices, where `N` is a type of the node label. Parsing with this combinator succeeds iff the vertex satisfies the predicate.
- `E[N](predicate: N =>Boolean)` the conbinator for processing edges, where `L` is a type of the edge label. Parsing with this combinator succeeds iff the edge satisfies the predicate.

To select cities which belong to some country, we can use the function `V: V[L]((e: Entity) =>e.country ="County␣Name")`. Here `Entity` is a property container for both edges and vertices. By using the `Dynamic` trait, all accesses to properties (like `(e: Entity).country`) are converted to the accesses to the properties of either a vertex or an edge. For the sake of simplicity, we will omit `Entity` type specifications for predicates. To query the graph for the paths from a city in the country X to a city in the country Y, we need to sequencially compose the combinators for selecting the appropriate cities. A sequentialital combinator `~` does just that: it sequentialy applies two queries one after the other. Let us denote a query for retrieving a city from the specific country `city(name: String)` and a query for retriving road edges `roadTo`. With this denotation, a query `city("X") ~ roadTo ~ city("Y")` returns the requested set of paths from the graph. The complete query with all needed subqueries is shown on fig. 5.

Having the query written, the next thing is to write a function to retrieve the actual data about the paths. If we care only about the names of the cities, we can return a pair of the cities for each path. First, we modify the query for vertices by adding the semantic action to it using the combinator `^`:
`def city(name: String) =syn(V(e.value() ==name) ^ (_.value))`.
Then we need to actually map a path to a pair of cities: this is done with the combinator `&`. The complete query is in fig. 6. The result is the sequence of pairs of cities with the road between them.

The whole set of basic combinators, which our library provides, is presented in the table 1. There are two kind of combinators the first kind combines parsers to form new parsers, meanwhile the second one is dedicated to the processing of the query result. Whenever a string is used withing a query, a parser which matches that string is implicitly generated.
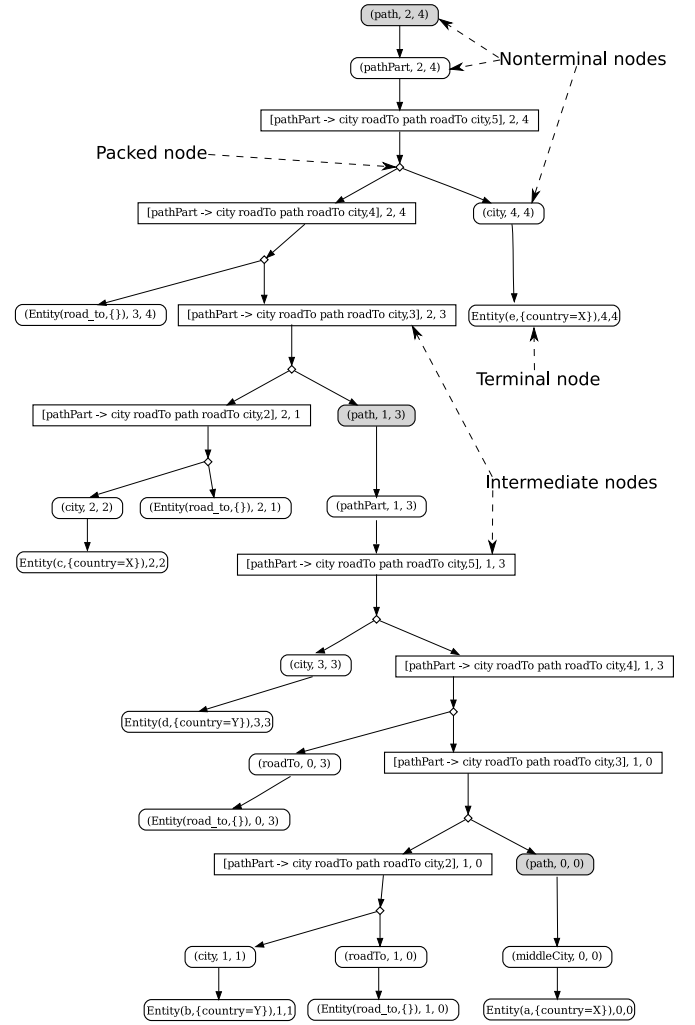
Figure 4: SPPF: result of applying cities query to the graph 3

```scala
def city(country: String) =
  V(e.country == country)
val roadTo = E(_.value() == "road_to")
val ourPath =
  city("X") ~ roadTo ~ city("Y")
```

Figure 5: Path query

4.2 Generic interface for input

The combinators in our library are independent of the input representation. It is enough to specify two basic combinators which handle vertices and edges. Vertex handling is checking wether the vertex satisfies the given predicate. In case of the edges, one needs to check which of the outgoing from the given vertex edges satisfy the given predicate. These two functions form the trait for the input. It has two type parameters: the type of edge labels `L` and the type of vertices labels `N`. We supported several different input sources:

```
def city(country: String) =
  syn(V(e.country() == country) ^ (_.name))
val roadTo = E(_.value() == "road_to")
val ourPath =
  syn(city("X") ~ roadTo ~ city("Y") &
    { case c0 ~ c1 => (c1, c2) }
```

Figure 6: Path query

Combinator	Description
a ~ b	sequential parsing: a then b
a | b	choice: a or b
a ?	optional parsing: a or nothing
a *	repetition of zero or more a
a +	repetition of at least one a
a ^ f	apply f function to a if a is a token
a ^^	capture output of a if a is a token
a & f	apply f function to a if a is a parser
a &&	capture output of a if a is a parser

Table 1: Meerkat combinators

- Neo4jInput — input source for the graph database Neo4J;
- GraphxInput — input source for the graph presented in memory using GraphX library;
- LinearInput — input source for the linear input data such as the ordinary strings.

```
trait Input[+L, +N] {
  def filterEdges(nodeId: Int,
      predicate: L => Boolean): Seq[(L, Int)]
  def checkNode(nodeId: Int,
      predicate: N => Boolean): Option[N]
}
```

Figure 7: Generalized input interface

Since the required functions are simple, we believe it is possible to support most storages of graph-structured data. Note, there are two technical limitations which arise from our decision to use values of the type Int as unique identifiers for vertices (the nodeId parameter). It is imposible to correctly query the input graph with more than MAX_INT vertices. Also there should be a way to provide such unique identifiers. Most systems already use unique identifiers by default, in other cases this feature can be implemented by the user.

4.3 Semantic Actions

Every path, which the query produces, has a derivation tree stored in the SPPF. The derivation tree is a very reach structure which can be hard to understand. To retrieve the actually useful for the user data, the library provides a mechanism of semantic actions. It is a way to apply some function to a parsed token or a subsequence.

Semantic action binder for the tokens—vertices and edges alike—is ^. The most common use for it is to extract properties from the token and combine them in some fashion.

```
// Defined in Terminal[+L] (edge) parser
def ^[U](f: L => U) =
  new SymbolWithAction[L, Nothing, U] {...


// Defined in Vertex[+N] parser
def ^[U](f: N => U) =
  new SymbolWithAction[Nothing, N, U] {...
```

For the combination of parsers there is a & binder. Being applied to a sequence of tokens, it can collect and process the data returned by the terminal parsers.

```
// Defined in Symbol[+L, +N, +V] parser
def &[U](f: V => U) =
  new SymbolWithAction[L, N, U] {...
```

They both produce a new parser that parses the input exactly like the given parser but also have a binded function. The function is referenced in each SPPF node produced by the corresponding parser.

The way semantic actions are executed has mostly remained the same as in the original Meerkat library: semantic actions are first executed for the children of the current node, collected and passed to a semantic action of the current node. If there are ambiguous nodes in the SPPF, the original Meerkat library just throws an exception. In our case, ambiguity can arise not only when there are multiple derivations of a string but ambiguous nodes can also represent several different subpaths which are derived from the respective nonterminal. We chose to provide a way to extract "all" derivation trees from the SPPF—lazily, since the number of the paths can be infinite. Unambiguous trees are yielded with a breadth-first search.

The composition of the extraction of trees and the semantic action execution is called executeQuery. It parses the input graph from all positions, produces a list of SPPF roots, extracts all derivations from every root, executes semantic actions and returns a lazy stream of results.

5 EXAMPLES

In this section, we describe some examples of queries written with the library proposed. We show that the combinatros are expressive enough for realistic queries and also ease their creation.

5.1 Complicated Query to Map

Let's query the city graph first. Let us capture one city, let's say city with name a. Now having a city graph and captured vertex we would like to know all paths such if as i city from begining of our path we visit country X then as i city from end of our path we visit country X too. And also the middle city in our path is our captured city a. In a terms of combinators we can define our path as shown on fig. 8. Here reduceChoice is a function which transforms a list of queries to one query which is formed by reducing given list with | combinator. The pathPart query recursively defines a path of our way. Also, middleCity is a vertex query which parses our captured city a and roadTo query parses a *roadTo* edge.

```
val countriesList = List("X", "Y")
val path =
  (reduceChoice(countriesList.map(pathPart)) |
    middleCity)
def pathPart(country: String) =
  syn(city(country) ~ roadTo ~ path ~
    roadTo ~ city(country))

val middleCity = V(_.value() == "a")
val roadTo = E(_.value() == "road_to")
def city(country: String) =
  V(_.country == country)
```

Figure 8: Path query

```
def reduceChoice(xs: List[Nonterminal]) =
  xs match {
    case x :: Nil    => x
    case x :: y :: xs =>
      syn(xs.foldLeft(x | y)(_ | _))
  }
```

Figure 9: Reduce choice function implementation

Now we would like, to get from our query only `city` combinator result. For that purpose let us modify it to make return result. In our library we have a `^` and `&` functions for that. Then we will have definition of our combinators as presented in fig. 10.

```
val middleCity =
  syn(V(_.value() == "a") ^^) & (List(_))
def pathPart(country: String) = syn(
  (city(country) ~ roadTo ~
    path ~ roadTo ~ city(country) & {
      case a ~ (b: List[_]) ~ Entity =>
        a +: b :+ c })
```

Figure 10: Fixed queries

Now we execute our query. It is evident that for the graph presented on fig. 3 we can get only three paths which satisfies given criteria:

- single-vertex path a;
- $b \rightarrow a \rightarrow d$
- $c \rightarrow b \rightarrow a \rightarrow d \rightarrow e$

A simplified SPPF for this query is presented in Fig. 4: rounded rectangles represent nonterminals and other rectangles represent productions. Every rectangle contains a nonterminal name or a production rule, as well as start and end nodes of the path in the input graph derived from the corresponding rectangle. Gray rectangles are start nonterminals.

5.2 Same Generation Query

Yet another example of first order functions usage is generalisation of classical same generation query which is one of basic context-free path queries. One of application of such queries is hierarchy analysing in RDF storages [35]. Let suppose that we have RDF graphs with two pairs of relation (each pair is relation and its revers): ($subClassOf$; $subClassOf^{-1}$) and ($type$; $type^{-1}$). We want to evaluate two queries which detect all pairs of nodes which are connected by path derivable in grammars G_1 (Fig. 11) and G_2 respectively (Fig. 12).

$$
\begin{aligned}
0: & \quad S \rightarrow subClassOf^{-1} \; S \; subClassOf \\
1: & \quad S \rightarrow type^{-1} \; S \; type \\
2: & \quad S \rightarrow subClassOf^{-1} \; subClassOf \\
3: & \quad S \rightarrow type^{-1} \; type
\end{aligned}
$$

Figure 11: Context-free grammar G_1 for query 1

$$
\begin{aligned}
0: & \quad S \rightarrow B \; subClassOf \\
0: & \quad S \rightarrow subClassOf \\
1: & \quad B \rightarrow subClassOf^{-1} \; B \; subClassOf \\
2: & \quad B \rightarrow subClassOf^{-1} \; subClassOf
\end{aligned}
$$

Figure 12: Context-free grammar G_2 for query 2

Of course, these queries can be written in Meerkat easely because it supports context-free queries: code is presented in Fig. 13 and Fig. 14.

```
val query1: Nonterminal = syn(
  "subclassof-1" ~ query1.? ~ "subclassof" |
  "type-1" ~ query1.? ~ "type")
```

Figure 13: The same generation query (Query 1) in Meerkat

```
val S = syn(
  "subclassof-1" ~ S ~ "subclassof")
val query2 = syn(S ~"subclassof")
```

Figure 14: The same generation query (Query 2) in Meerkat

As you can see, grammars and code representations for these two queries looks pretty similar. May we avoid code duplication and generalize them? Yes, we can and not only for these two queries. The function `sameGen` presented in Fig 15 is a generalization of the same generation query and is independent of the environment such as the input graph structure or other parsers and also uses a function `reduceChoice` presented in 9. It can be used for the creation of other queries, including the one presented in Fig 13: it is the result of the application of `sameGen` to the appropriate relations (which can be treated as opening and closing brackets). Another application of the `sameGen` is a Query 2, which can be founded in Fig. 17.

```
def sameGen(brs) =
  reduceChoice(
    bs.map { case (lbr, rbr) =>
      lbr ~ syn(sameGen(bs).?) ~ rbr })
```

Figure 15: Generic function for the same generations query

```
val query1 = syn(sameGen(List(
  ("subclassof-1", "subclassof"),
  ("type-1", "type"))))
```

Figure 16: Query 1 as an application of `sameGen`

```
val query2 = syn(
  sameGen(List(("subclassof-1", "subclassof"))) ~
    "subclassof")
```

Figure 17: Query 2 as an application of `sameGen`

We show that parser combinators provide a simple and safe way to creation of generic queries. By using this ability, it may be possible to create a library of "standard templates" for most popular generic queries like same generation query or for domain specific queries (for example, for specific static code analysis problem).

5.3 Classical Movies Queries

In order to demonstarte expressive power of our solution and to demonstrate more scenarios for semantic actions usage we preovide some examples of classical queries to movie database which represents movies, actors, directors, users and relationships between them. All of queries can be found on a Neo4j tutorial page [5].

First of all, we need to introduce some helpers for nodes and edges porcessing simplification. They are built upon basic combinators from our library and presented in figure 18.

```
def LV(labels: String*) = V((e: Entity) => labels.forall(e.hasLabel))
def outLE(label: String) = outE((e: Entity) => e.label() == label)
def inLE(label: String) = inE((e: Entity) => e.label() == label)
```

Figure 18: Helpers for edges and nodes processing

Let's look at one of these examples written using Cypher language and how it transforms into Meerkat query (fig. 20).

Firstly, we should define parsers that correspond to persons, movies and specific user. Secondly, we should also create parsers corresponding to relations. Then, we can get a full path parser using combination of previously defined primitives. Finally, adding semantic action that extracts all needed data we get a complete path query. Result is shown on fig. 21. Using similar transformations we can write the most queries which can be defined using Cypher.

We show that our library is expressive enough to fromulate realistic queries. Also we demonstrate some cases of semantics actions usage. Main difference from Cypher is that library provides only path querying mechanusms, so all additional logic such as filtering, sorting or grouping must be implemented manually as a separated step.

Also, we detect that our library doesn't support incoming edges precessing which may be useful in some scenarios.

6 EVALUATION

In this section, we present an evaluation of our graph querying library. We measure its performance on a classical set of ontology graphs [35]: both when the graph is loaded into the RAM and for the integration with the Neo4j database and compare it with the solution based on the GLL parsing algorithm [6] and the Trails [12] library for graph traversals. We also show how may-alias static code analysis can be done by the means of the library.

All tests have been performed on a computer running Fedora 27 with quad-core Intel Core i7 2.5 GHz CPU and 8 GB of RAM.

6.1 Ontology querying

Querying for ontologies is a well-known graph querying problem. We evaluate our library on some popular ontologies in form of RDF files from the paper [35]. First, we convert RDF files to a labelled directed graph in the following manner: we create two edges $(subject, predicate, object)$ and $(object, predicate^{-1}, subject)$ for every RDF triple $(subject, predicate, object)$. Then the graph is either loaded into the Neo4j database or is loaded directly into the memory. Then we run two queries from the paper [6] for these graphs. The grammars for the queries are presented in Fig. 16 and Fig. 17.

The performance results are shown in the table 2 where $#results$ is a number of pairs of the nodes between which exists at least one S-path.

The Meerkat-based and the GLL-based [6] solutions show the same results (column $#results$) and the Query 1 runs up to two times faster on the Meerkat-based solution than on the GLL-based, meanwhile the GLL-based solution is faster for the Query 2. Querying the database is naturally $2 - 4$ times slower than querying the graphs located in the RAM.

6.2 Static code analysis

Alias analysis is a fundamental static analysis problem [15]: it checks may-alias relations between code expressions and can be formulated as a context-free language (CFL) reachability problem [23], which is closely related to the context-free path querying problem. In this analysis, a program is represented as a Program Expression Graph (PEG) [36]. Vertices in a PEG correspond to program expressions and edges are relations between them. There are two types of edges possible while analyzing C source code: **D**-edge and **A**-edge.

- Pointer dereference edge (**D**-edge). For each pointer dereference $*e$ there is a directed **D**-edge from e to $*e$.
- Pointer assignment edge (**A**-edge). For each assignment $*e_1 = e_2$ there is a directed **A**-edge from e_2 to $*e_1$

For the sake of simplicity, we add edges labelled by $\overline{D}$ and $\overline{A}$ which correspond to the reversed **D**-edge and **A**-edge, respectively.

Ontology	#nodes	#edges	Query 1					Query 2				
			#results	In memory graph (ms)	DB query (ms)	Trails (ms)	GLL (ms)	#results	In memory graph (ms)	DB query (ms)	Trails (ms)	GLL (ms)
atom-primitive	291	685	15454	64	67	2849	232	122	29	79	453	19
biomedical-measure-primitive	341	711	15156	112	108	3715	482	2871	18	18	60	26
foaf	256	815	4118	11	11	432	29	10	1	1	1	1
funding	778	1480	17634	69	68	367	179	1158	9	9	76	13
generations	129	351	2164	4	4	9	12	0	1	0	0	0
people_pets	337	834	9472	37	37	75	80	37	1	1	2	1
pizza	671	2604	56195	333	325	7764	793	1262	17	18	905	50
skos	144	323	810	1	2	6	6	1	1	0	0	0
travel	131	397	2499	11	13	34	21	63	1	1	1	2
univ-bench	179	413	2540	10	10	31	24	81	1	1	2	1
wine	733	2450	66572	405	401	3156	606	133	2	3	4	5

Table 2: Comparation of Meerkat, Trails and GLL performance on ontologies

The grammar for may-alias problem from [36] is presented in fig. 22. It contains two nonterminals **M** and **V**, which we can query for.

- **M**-production means that two l-value expressions are memory aliases i.e. may refer to the same memory location.
- **V**-production means that two expressions are value aliases i.e. may evaluate to the same pointer value.

We run the **M** and **V** queries on some open-source C projects: the results are in the table 3. We conclude that our solution is expressive and performant enough to be used for static analysis problems, which can be expressed as CFPQs.

6.3   Classical Movies Queries

In order to examine the applicability of our library for the real data processing we evaluate queries presented in the section 5 on the classical movie database which contains more than 60000 nodes and !!! edges.

All queries are performed using the Neo4j database connected to Meerkat. During evaluation of these queries, caching in Neo4j was disabled to simplify time measurement.

Our library implements a general parsing algorithm which supports abitrary context-free queries. Unfortunately, when the queries are regular, there is a big overhead for processing them. The queries mentioned in this section are all regular, thus it is not a surprise that our implementation is dramatically slower that the native Neo4j solution. Besides the overhead itself, our library does not take any advantages from the internal Neo4j optimizations.

Despite the seeming failure of this test, we still can conclude that the combinators can process quite big datasets in a reasonable time: 6 times bigger in terms of the number of vertices than the other tests. It also shows that further optimization of the realization is justified.

6.4   Comparison with Trails

Trails [12] is a Scala graph combinator library. It provides traversers for describing paths in graphs which resemble parser combinators

and calculates possibly infinite stream of all possible paths described by the composition of the basic traversers. Both Trails and Meerkat-based solution support parsing of the graphs located in RAM, so we compare the performance of the Trails and the Meerkat-based solution on the ontology queries described above: the results are presented in the table 2. Trails and Meerkat-based solution compute the same results, but Trails is up to 10 times slower than the Meerkat-based solution.

To summarize, we demonstrated that parser combinators suit for implementing real queries. Our implementation is as performant as the other existing combinators library and is comparable to the GLL-based solution.

7   DISCUSSION AND CONCLUSION

We propose a native way to integrate a language for context-free path querying into a general-purpose programming language. Our solution handles arbitrary context-free grammars and arbitrary input graphs. The proposed approach is language-independent and may be implemented in nearly every general-purpose programming language. We implement it in Scala and show that our approach can be applied to the real world problems

We can propose some possible directions for the future work.

First of all, in order to improve the performance and investigate the scalability of the solution, we plan to implement a parallel single machine and distributed GLL. It is a challenge from both the theoretical and the practical standpoint. (!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!! This needs more thought !!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!)

Another direction is improvements of semantic actions computation, otherwise known as the attribute grammars handling. Although it is impossible in general, techniques such as lazy evaluation can provide a technically adequate solution which is demonstarted in our solution. But our implmentation is very naive and should be improved. One of possible direction is utilizatuin of relational programming (minikanren) which is aimed to search [?]Ekaterina, we need your help!!!. For what class of semantic actions it is possible to provide a precise general solution is a theoretical question to be answered.

Program	#edges	#nodes	Code Size (KLOC)	In memory graph (ms)	Neo4j graph (ms)	Trails graph (ms)	M aliases	V aliases
wc-5.0	332	770	0.5	0	2	3	174	107
pr-5.0	815	2062	1.7	11	12	14	1131	63
ls-5.0	1687	4734	2.8	43	51	170	5682	253
bzip2-1.0.6	632	1508	5.8	8	13	21	813	71
gzip-1.8	2687	7510	31	111	120	537	4567	227

Table 3: Running may-alias queries on Meerkat on some C open-source projects

Query	Neo4j + Meerkat (ms)	Neo4j + Cypher (ms)
Mutual Friend recommendations	120	16
Directed more than 2 films, Acted in more than 10	423	30
Find the most profilic actors	4528	222
Actors who played in some film	680	9

Table 4: Running regular queries using Meerkat and Cypher

Some important problems in static code analysis reqwuire languages more expressive than contex-free one. For example, context-sensitive data-dependence analysis may be precisely expressed in terms of linear-conjunctive language [19] reachability, but not context-free [34]. While problem formulation is precise, it is possible to get only approximated solution, because emptyness problem for linear-conjunctive languages is undersidable. It would be an interesting task to support not only linear-cinjunctive grammars, but arbitrary conjunctive grammars [18] in te library and investigate nature of approximation. Finally it would be interesting to create a core for static analysis framework based on language reachcbility and compare it with the respective part of Flix or other related tools.

Improved version of OpenCypher [14], which is the one of the most popular graph query languages, provides context-free path querying mechanism. Detailed comparison with it may provide more information for direcion of future work.

ACKNOWLEDGMENTS

REFERENCES

[1] Pablo Barceló, Gaelle Fontaine, and Anthony Widjaja Lin. 2013. Expressive Path Queries on Graphs with Data. In *International Conference on Logic for Programming Artificial Intelligence and Reasoning*. Springer, 71–85.

[2] Osbert Bastani, Saswat Anand, and Alex Aiken. 2015. Specification inference using context-free language reachability. In *ACM SIGPLAN Notices*, Vol. 50. ACM, 553–566.

[3] Tevfik Bultan, Fang Yu, Muath Alkhalaf, and Abdulbaki Aydin. 2018. String Analysis for Software Verification and Security. (2018).

[4] James Cheney, Sam Lindley, and Philip Wadler. 2013. A Practical Theory of Language-integrated Query. *SIGPLAN Not.* 48, 9 (Sept. 2013), 403–416. https://doi.org/10.1145/2544174.2500586

[5] Andrei Marian Dan, Manu Sridharan, Satish Chandra, Jean-Baptiste Jeannin, and Martin Vechev. 2017. Finding Fix Locations for CFL-Reachability Analyses via Minimum Cuts. In *International Conference on Computer Aided Verification*. Springer, 521–541.

[6] Semyon Grigorev and Anastasiya Ragozina. 2017. Context-free Path Querying with Structural Representation of Result. In *Proceedings of the 13th Central & Eastern European Software Engineering Conference in Russia (CEE-SECR '17)*. ACM, New York, NY, USA, Article 10, 7 pages. https://doi.org/10.1145/3166094.3166104

[7] Jelle Hellings. 2014. Conjunctive context-free path queries. (2014).

[8] Jelle Hellings. 2015. Path Results for Context-free Grammar Queries on Graphs. *CoRR* abs/1502.02242 (2015). http://arxiv.org/abs/1502.02242

[9] Wei Huang, Yao Dong, Ana Milanova, and Julian Dolby. 2015. Scalable and precise taint analysis for android. In *Proceedings of the 2015 International Symposium on Software Testing and Analysis*. ACM, 106–117.

[10] Anastasia Izmaylova, Ali Afroozeh, and Tijs van der Storm. 2016. Practical, General Parser Combinators. In *Proceedings of the 2016 ACM SIGPLAN Workshop on Partial Evaluation and Program Manipulation (PEPM '16)*. ACM, New York, NY, USA, 1–12. https://doi.org/10.1145/2847538.2847539

[11] Mark Johnson. 1995. Memoization in Top-down Parsing. *Comput. Linguist.* 21, 3 (Sept. 1995), 405–417. http://dl.acm.org/citation.cfm?id=216261.216269

[12] Daniel Kröni and Raphael Schweizer. 2013. Parsing Graphs: Applying Parser Combinators to Graph Traversals. In *Proceedings of the 4th Workshop on Scala (SCALA '13)*. ACM, New York, NY, USA, Article 7, 4 pages. https://doi.org/10.1145/2489837.2489844

[13] Kazumasa Kumamoto, Toshiyuki Amagasa, and Hiroyuki Kitagawa. 2015. A System for Querying RDF Data Using LINQ. In *Network-Based Information Systems (NBiS), 2015 18th International Conference on*. IEEE, 452–457.

[14] Tobias Lindaaker. 2017. OpenCypher Path Patterns (CIP2017-02-06 Path Patterns). (2017). https://github.com/thobe/openCypher/blob/rpq/cip/1.accepted/CIP2017-02-06-Path-Patterns.adoc#153-compared-to-context-free-languages

[15] Thomas J. Marlowe, William G. Landi, Barbara G. Ryder, Jong-Deok Choi, Michael G. Burke, and Paul Carini. 1993. Pointer-induced Aliasing: A Clarification. *SIGPLAN Not.* 28, 9 (Sept. 1993), 67–70. https://doi.org/10.1145/165364.165387

[16] Erik Meijer, Brian Beckman, and Gavin Bierman. 2006. LINQ: Reconciling Object, Relations and XML in the .NET Framework. In *Proceedings of the 2006 ACM SIGMOD International Conference on Management of Data (SIGMOD '06)*. ACM, New York, NY, USA, 706–706. https://doi.org/10.1145/1142473.1142552

[17] A. Mendelzon and P. Wood. 1995. Finding Regular Simple Paths in Graph Databases. *SIAM J. Computing* 24, 6 (1995), 1235–1258.

[18] Alexander Okhotin. 2001. Conjunctive grammars. *Journal of Automata, Languages and Combinatorics* 6, 4 (2001), 519–535.

[19] Alexander Okhotin. 2003. On the Closure Properties of Linear Conjunctive Languages. *Theor. Comput. Sci.* 299, 1 (April 2003), 663–685. https://doi.org/10.1016/S0304-3975(02)00543-1

[20] Polyvios Pratikakis, Jeffrey S Foster, and Michael Hicks. 2006. Existential label flow inference via CFL reachability. In *SAS*, Vol. 6. Springer, 88–106.

[21] Eric Prud, Andy Seaborne, et al. 2006. SPARQL query language for RDF. (2006).

[22] Joan Gerard Rekers. 1992. *Parser generation for interactive environments*. Ph.D. Dissertation. Universiteit van Amsterdam.

[23] Thomas Reps. 1997. Program Analysis via Graph Reachability. In *Proceedings of the 1997 International Symposium on Logic Programming (ILPS '97)*. MIT Press, Cambridge, MA, USA, 5–19. http://dl.acm.org/citation.cfm?id=271338.271343

[24] Thomas Reps, Susan Horwitz, and Mooly Sagiv. 1995. Precise Interprocedural Dataflow Analysis via Graph Reachability. In *Proceedings of the 22Nd ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages (POPL '95)*. ACM, New York, NY, USA, 49–61. https://doi.org/10.1145/199448.199462

[25] Juan L Reutter, Miguel Romero, and Moshe Y Vardi. 2015. Regular queries on graph databases. *Theory of Computing Systems* (2015), 1–53.

[26] Marko A Rodriguez. 2015. The gremlin graph traversal machine and language (invited talk). In *Proceedings of the 15th Symposium on Database Programming Languages*. ACM, 1–10.

[27] Elizabeth Scott and Adrian Johnstone. 2010. GLL parsing. *Electronic Notes in Theoretical Computer Science* 253, 7 (2010), 177–189.
[28] Elizabeth Scott and Adrian Johnstone. 2013. GLL parse-tree generation. *Science of Computer Programming* 78, 10 (2013), 1828–1844.
[29] Elizabeth Scott, Adrian Johnstone, and Rob Economopoulos. 2007. BRNGLR: a cubic Tomita-style GLR parsing algorithm. *Acta informatica* 44, 6 (2007), 427–461.
[30] Petteri Sevon and Lauri Eronen. 2008. Subgraph queries by context-free grammars. *Journal of Integrative Bioinformatics* 5, 2 (2008), 100.
[31] Kai Wang, Aftab Hussain, Zhiqiang Zuo, Guoqing Xu, and Ardalan Amiri Sani. 2017. Graspan: A Single-machine Disk-based Graph System for Interprocedural Static Analyses of Large-scale Systems Code. In *Proceedings of the Twenty-Second International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS '17)*. ACM, New York, NY, USA, 389–404. https://doi.org/10.1145/3037697.3037744
[32] Dacong Yan, Guoqing Xu, and Atanas Rountev. 2011. Demand-driven Context-sensitive Alias Analysis for Java. In *Proceedings of the 2011 International Symposium on Software Testing and Analysis (ISSTA '11)*. ACM, New York, NY, USA, 155–165. https://doi.org/10.1145/2001420.2001440
[33] Mihalis Yannakakis. 1990. Graph-theoretic methods in database theory. In *Proceedings of the ninth ACM SIGACT-SIGMOD-SIGART symposium on Principles of database systems*. ACM, 230–242.
[34] Qirun Zhang and Zhendong Su. 2017. Context-sensitive Data-dependence Analysis via Linear Conjunctive Language Reachability. In *Proceedings of the 44th ACM SIGPLAN Symposium on Principles of Programming Languages (POPL 2017)*. ACM, New York, NY, USA, 344–358. https://doi.org/10.1145/3009837.3009848
[35] Xiaowang Zhang, Zhiyong Feng, Xin Wang, Guozheng Rao, and Wenrui Wu. 2016. Context-free path queries on RDF graphs. In *International Semantic Web Conference*. Springer, 632–648.
[36] Xin Zheng and Radu Rugina. 2008. Demand-driven Alias Analysis for C. In *Proceedings of the 35th Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages (POPL '08)*. ACM, New York, NY, USA, 197–208. https://doi.org/10.1145/1328438.1328464

```scala
val query = syn((LV("Movie") :: V((_:Entity).title == "Forrest
                syn(LV("Actor") ^ (_.getProperty[String]

    executeQuery(query, input).foreach(println)



    =======================
val query = syn((syn(LV("Actor") ^^) ~ outLE("ACTS_IN") ~ LV("
    ((a: Entity) => (a.getProperty[String]("name"), a.id.as

    executeQuery(query, input)
      .groupBy({case (a, i) => i})
      .toIndexedSeq
      .map({case (i, ms) => (ms.head._1, ms.length)})
      .sortBy({case (a, mc) => -mc})
      .take(10)
      .foreach({case (a, mc) => println(a, mc)})


========================

val directors = syn((syn(LV("Actor", "Director") ^^) ~ outLE("
                     & (d
=> d.id.asInstanceOf[String].toInt))

    val directorsMap = executeQuery(directors, input)
      .groupBy(i => i)
      .map({case (i, ms) => (i, ms.length)})
      .filter({case (_, ms) => ms >= 2})

    val actor_prof_director = syn(LV("Actor", "Director") :: V

    val acts = syn((actor_prof_director ~ outLE("ACTS_IN") ~ L
      (a => (a.getProperty[String]("name"), a.id.asInstanceOf[

    executeQuery(acts, input)
      .groupBy({case (a, i) => i})
      .toStream
      .map({case (i, ms) => (i, ms.head._1, ms.length)})
      .filter({case (i, a, mc) => mc >= 10})
      .foreach({case (i, a, mc) => println((a, mc, directorsMa

==========================================

    val adilfulara = syn(LV("User") :: V((_: Entity).login ==

    val query = syn((adilfulara ~ (outLE("FRIEND") | inLE("Fri
                                  syn(outLE("RATED") ^^) ~ syn(
      {case p ~ r ~ m => (p.getProperty[String]("name"), m.tit
                          if (r.hasProperty("comment")) r.comm

    executeQuery(query, input)
      .filter({case (_, _, s, _) => s > 3})
      .foreach(println)


      =====================================

      MATCH (m:Movie {title: 'Forrest_Gump'})<-[:ACTS_IN]-(a:A

RETURN a.name, a.birthplace
====================
```

```
MATCH (u:User {login: 'adilfulara'})-[:FRIEND]->
      (f:Person)-[r:RATED]->(m:Movie)
WHERE r.stars > 3
RETURN f.name, m.title, r.stars, r.comment
```

Figure 20: Mutual Friend recommendations query in Cypher

```
val adilfulara =
    syn(V((e: Entity) => e.hasLabel("User")
                && e.login == "adilfulara"))
val friend =
    syn(E((e: Entity) => e.value() == "FRIEND"))

val person =
    syn(V((e: Entity) => e.hasLabel("Person")) ^^)

val rated =
    syn(E((e: Entity) => e.value() == "RATED") ^^)

val movie =
    syn(V((e: Entity) => e.hasLabel("Movie")) ^^)

val query =
    syn((adilfulara ~ friend ~ person ~
         rated ~ movie) &
    {case p ~ r ~ m =>
        (p.name, m.title, r.stars.asInstanceOf[Int],
         if (r.hasProperty("comment")) r.comment
         else "")})

executeQuery(query, input)
  .filter({case (_, _, s, _) => s > 3})
```

Figure 21: Mutual Friend recommendations (Meerkat query)

$$M \rightarrow \overline{D}\, V\, D$$
$$V \rightarrow (M?\, \overline{A})^*\, M?\, (A\, M?)^*$$

Figure 22: Context-Free grammar for the may-alias problem

```
val M = syn("nd" ~ V ~ "d")
val V = syn((M.? ~ "na").* ~ M.? ~ ("a" ~ M.?).*)
```

Figure 23: Meerkat representation of may-alias problem grammar