Parser-Combinators for Context-Free Path Querying

Sophia Smolina
Electrotechnical University
St. Petersburg, Russia
sofysmol@gmail.com

Ekaterina Verbitskaia
Saint Petersburg State University
St. Petersburg, Russia
kajigor@gmail.com

Ilya Kirillov
Saint Petersburg State University
St. Petersburg, Russia
kirillov.ilija@gmail.com

Semyon Grigorev
Saint Petersburg State University
St. Petersburg, Russia
semen.grigorev@jetbrains.com

ABSTRACT

A transparent integration of a domain-specific language for specification of context-free path queries (CFPQs) into a general-purpose programming language as well as static checking of errors in queries may greatly simplify the development of applications utilizing CFPQs. Such techniques as LINQ and ORM can be used for the integration, but they have issues with flexibility: query decomposition and reusing of subqueries are a challenge. Adaptation of parser combinators technique for paths querying may solve these problems. Conventional parser combinators process linear input and only the Trails library is known to apply this technique for path querying. Trails suffers the common parser combinators issue: it does not support left-recursive grammars and also experiences problems in cycles handling. We demonstrate that it is possible to create general parser combinators for CFPQ which support arbitrary context-free grammars and arbitrary input graphs. We implement a library of such parser combinators and show that it is applicable for realistic tasks.

CCS CONCEPTS

• **Information systems → Graph-based database models**; **Query languages for non-relational engines**; • **Software and its engineering** → *Functional languages*; • **Theory of computation** → *Grammars and context-free languages*;

KEYWORDS

Graph Databases, Language-Constrained Path Problem, Context-Free Path Querying, Parser Combinators, Generalized LL, GLL, Neo4J, Scala

1 INTRODUCTION

One useful type of graph queries is language-constrained path queries [3]. There are several languages for graph traversing/querying which support constraints formulated in terms of regular languages. For example SPARQL [23], Cypher [1], and Gremlin [28]. In this work, we are focused on context-free path queries (CFPQs) which use context-free languages for constraints specification and are used in bioinformatics [31], static code analysis [4, 22, 25, 33, 37], and RDF processing [36]. There are a lot of problem-specific solutions and theoretical research on CFPQs [2, 7–9, 19, 27, 31, 34]. Among them, cfSPARQL [36] is a single known graph query language to support CF constraints. Generic solution for the integration of CFPQs into general-purpose languages is not discussed enough.

When one develops a data-centric application, one wants to use a general purpose programming language and have a transparent and native access to data sources. String-embedded DSLs is one way to do it. It utilizes a driver to execute a query written as a string and to return a possibly untyped result. This approach has serious drawbacks. First of all, a DSL may require additional knowledge from a developer. Moreover, a string-embedded language itself is a source of possible errors and vulnerabilities, static detection of which is very difficult [5]. In trying to solve these issues, such special techniques as Object Relationship Mapping (ORM) or Language Integrated Query (LINQ) [6, 14, 18] were created. Unfortunately, they still experience difficulties with flexibility: for example with the query decomposition and the reusing of subqueries. In this paper, we propose a transparent and natural integration of CFPQs into a general-purpose language.

It is necessary to find an appropriate technique for integration of context-free language specification into general-purpose programming languages. One natural way to specify a language is to specify its formal grammar which can be done by using a special DSL based, for example, on EBNF-like notation [32]. The classical alternative way is parser combinators [10] which provide all required features, including transparent integration, compile-time checks of correctness, high-level techniques for generalization.

An idea to use combinators for graph traversing has already been proposed in [13], but the solution presented provides only approximated handling of cycles in the input graph and does not support left-recursive grammars. Authors pointed out that the idea

---

[1]Cypher language web page: https://neo4j.com/developer/cypher-query-language/. Access date: 16.01.2018

described is very similar to the classical parser combinators, but the language class supported or restrictions are not discussed. This point is very important, because conventional combinators implement top-down parsing and cannot handle left-recursive and ambiguous grammars.

In [11], authors demonstrate a set of parser combinators which can handle arbitrary context-free grammars by using ideas of Generalized LL [29] (GLL). Meerkat [2] parser combinators library is based on [11] and provides result of parsing in a compact form as Shared Packed Parse Forest [24] (SPPF). It is showed that SPPF is a suitable finite structural representation of a CFPQ query result, even if the set of paths is infinite [7], which can be used for paths extraction, queries debugging and processing of result.

In this paper, we show how to compose these ideas and present the parser combinators for CFPQ which can handle arbitrary context-free grammars and provide a structural representation of the result. We make the following contributions in this paper.

(1) We show that it is possible to create parser combinators for context-free path querying which work on both arbitrary context-free grammars and arbitrary graphs and provide a finite structural representation of the query result.

(2) We provide the implementation of the parser combinators library in Scala. This library provides an integration to Neo4j [3] graph database. The source code is available on GitHub:https://github.com/YaccConstructor/Meerkat.

(3) We perform an evaluation on realistic data. Also, we compare the performance of our library with another GLL-based CFPQ tool and with the Trails library. We conclude that our solution is expressive and performant enough to be applied to the real-world problems.

2 GENERALIZED LL

Sott, Ali, Meerkat

handle arbitrary context-free grammar, cubic time complexity. GLL-based combinators.

Meerkat library is a general parser combinators library; by using memoization, continuation-passing style and the ideas of Johnson [12], it supports arbitrary context-free specifications.

2.1 SPPF

Structural representation. Derevation tree. Forest for unambiguous grammars. For graph too. Shared Packed Oarse Forest (SPPF) [24] structure, description, usability for CFPQ.

Binarized Shared Packed Parse Forest (SPPF) [30] compresses derivation trees optimally reusing common nodes and subtrees. Version of GLL which utilizes this structure for parsing forest representation achieves worst-case cubic space complexity [?].

Binarized SPPF can be represented as a graph in which each node has one of four types described below. We denote the start and the end positions of substring as i and j respectively, and we call tuple (i, j) an *extension* of a node.

- **Terminal node** with label (i, T, j).

- **Nonterminal node** with label (i, N, j). This node denotes that there is at least one derivation for substring $\alpha = \omega[i..j-1]$ such that $N \Rightarrow_G^* \alpha, \alpha = \omega[i..j-1]$. All derivation trees for the given substring and nonterminal can be extracted from SPPF by left-to-right top-down graph traversal started from respective node.

- **Intermediate node**: a special kind of node used for binarization of SPPF. These nodes are labeled with (i, t, j), where t is a grammar slot.

- **Packed node** with label $(N \to \alpha, k)$. Subgraph with "root" in such node is one possible derivation from nonterminal N in case when the parent is a nonterminal node labeled with $(\diamondsuit (i, N, j))$.

An example of SPPF is presented in figure **??**. We remove redundant intermediate and packed nodes from the SPPF to simplify it and to decrease the size of the structure.

2.2 GLL for CFPQ

Our work [7]. It is possible to use GLL for CFPQ. String-embedded querying is a ugly solution.

As the Meerkat library is closely related to the Generalized LL algorithm and since GLL can be generalized for context-free path querying [7], it is also possible to adapt Meerkat library for graph querying. It can be done by providing a function for retrieving the symbols which follow the specified position and utilizing it in the basic set of combinators. Detail described below.

3 PARSER COMBINATORS FOR PATH QUERYING

Parser combinators provide a way to specify a language syntax in terms of functions and operations on them. A parser in this framework is usually a function which consumes a prefix of an input and returns either a parsing result or an error if the input is erroneous. Parsers can be composed by using a set of parser combinators to form more complex parsers. A parser combinators library provides with a set of basic combinators (such as sequential application or choice), and there can also be user-defined combinators. Most parser combinators libraries, including the Meerkat library, can only process the linear input — strings or some kind of streams. We extend the Meerkat library to work on the graph input.

3.1 The set of combinators

For creating queries we need to work with edges and vertices. There are two main combinators for that:

- V[L](predicate: L =>Boolean) combinator for working with vertices. Accepts a predicate and parses only vertices which satisfies that predicate

- E[N](predicate: N =>Boolean) combinator for working with edges. Accepts a predicate and parses only edges which satisfies that predicate

(Write sth about syn macro)

Let's assume we have a city graph. In that graph vertices are cities. There can be a road from one city to another, this relation is shown in graph as an edge with label *road_to*. Every city belongs to a country. Such graph is shown on 7. Now we would like to get all
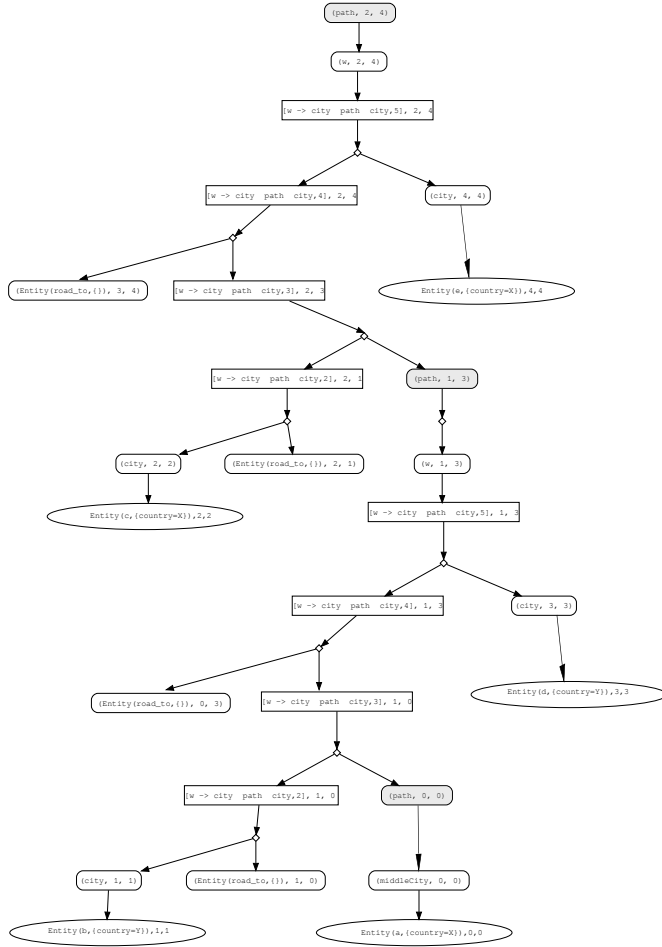
Figure 1: SPPF: result of applying actor/movie query to the graph 7

cities from our graph which belongs to some country. For that we will write V[L]((e: Entity) =>e.country ="County_Name"). Here Entity is a property container for graph entities: edges and vertices. Also, for the sake of simplicity, we will not explicitly specify Entity type for predicates. Now let us build a query which gets all roads from city in *country*0 to city in *country*. For that we can use a sequentialital combinator ~. It allows to create queries which sequentialy applies two queries one after another. When we have subquery for retrieving a vertex with specific city, let's call it city(name: String) and a subquery roadTo for retrieving road edges. Let's finally build a query city("city0") ~ roadTo ~ city("city1"). The full query with subqueries is shown on Fig. 2.

```
def city(name: String) =
  V(e.name == name)
val roadTo = E(_.value() == "road_to")
val ourPath = city("country0") ~ roadTo ~ city("country1")
```

Figure 2: Path query

Now we would like to get all pair of cities which have a road between them. So we need to transform our query to use semantic actions which is described in 3.4 section. Now let us specify what we want from every our query. From the city query we want only city name, so we need to map a result of basic vertex combinator. For that case we have a ^ combinator we can write
def city(name: String) =syn(V(e.value() ==name) ^ (_.value))
to achive that. In ourPath query we need first and second cities to be represented as a pair. For that we have a & combinator which will map our sequence to a pair of strings. The final representation is shown on 3. Now then we run that query we will get a list which consists of all pairs of city's names which have a road between.

```
def city(country: String) =
  syn(V(e.country() == country) ^ (_.name))
val roadTo = E(_.value() == "road_to")
val ourPath =
  syn(city("city0") ~ roadTo ~ city("city1") &
    {case c0 ~ c1 => (c1, c2)}
```

Figure 3: Path query

The whole set of basic combinators our library provides are presented in table 1. It consists of two kind of combinators. The first kind creates new parsers from existing ones, meanwhile the second one allows mapping parsers result. Parsers for matching strings are implicitly generated whenever a string is used within a query. The classical same generation query [1] can be written using the library as presented in Fig. 10.

Combinator	Description
a ~ b	sequential parsing: a then b
a | b	choice: a or b
a ?	optional parsing: a or nothing
a *	repetition of zero or more a
a +	repetition of at least one a
a ^ f	apply f function to a if a is a token
a ^^	capture output of a if a is a token
a & f	apply f function to a if a is a parser
a &&	capture output of a if a is a parser

Table 1: Meerkat combinators

3.2 More complicated example

Let's form a complex query for our city graph. Let us capture one city, let's say city *a*. Now having a city graph and captured vertex

we would like to know all paths such if as *i* city from begining of our path we visit country *X* then as *i* city from end of our path we visit country *X* too. And also the middle city in our path is our captured city *a*. In a terms of combinators we can define our path as shown on Fig. 4. Here `reduceAsOr` is a function which transforms a list of queries to one query which is formed by reducing given list with | combinator. The `pathPart` query recursively defines a path of our way. Also, `middleCity` is a vertex query which parses our captured city *a* and `roadTo` query parses a roadTo edge.

```
val countriesList = List("X", "Y")
val path =
  (reduceAsOr(countriesList.map(pathPart)) | middleCity)
def pathPart(country: String) =
  syn(city(country) ~ roadTo ~ path ~ roadTo ~ city(country))

val middleCity = V(_.value() == "a")
val roadTo = E(_.value() == "road_to")
def city(country: String) =
  V(_.country == country)
```

Figure 4: Path query

The most exciting feature of our library is that queries can be used as first-class values which means greater generalization and composition. The function `reduceAsOr` presented in Fig 5 is a generalization of bracket sequence query and is independent of the environment such as the input graph structure or other parsers. It can be used for the creation of other queries, including the one presented in Fig 4: it is the result of the application of `reduceAsOr` to the appropriate relations (which can be treated as opening and closing brackets).

```
def reduceAsOr(xs: List[Nonterminal]) =
  xs match {
    case x :: Nil    => x
    case x :: y :: xs =>
      syn(xs.foldLeft(x | y)(_ | _))
  }
```

Figure 5: Combinators implementation

Now we would like, to get from our query only `city` combinator result. For that purpose let us modify it to make return result. In out library we have a ^ and & functions for that. Then we will have the following definition of our combinators:

Now we execute our query. For the graph presented on fig. 7 we will get paths which satisfies given criteria a, b->a->d, c->b->a->d->e

A simplified SPPF for this query is presented in Fig. 1: rounded rectangles represent nonterminals and other rectangles represent productions. Every rectangle contains a nonterminal name or a production rule, as well as start and end nodes of the path in the input graph derived from the corresponding rectangle. Gray rectangles are start nonterminals.

```
val middleCity =
  syn(V(_.value() == "a") ^^) & (List(_))
def pathPart(country: String) = syn(
  (city(country) ~ roadTo ~
    path ~ roadTo ~ city(country) & {
      case a ~ (b: List[_]) ~ Entity =>
        a +: b :+ c })
```

Figure 6: Fixed combinators

Figure 7: Example Input graph

3.3 Generic interface for input

Combinators is a way to describe a query. When we have a query we may want to execute that query on some graph considering it as an input for our query. The cool thing is that that input can be generalized to achive flexibility of all possible application of graph combinators and support a wide range of data sources. The idea is to present the input as two functions, one for working with edges and one for vertices. The first one would allow to get all edges outcoming from current vertex and also satisfies given predicate. The second one will allow to check if current vertex satisfies given predicate. That interface is presented on Fig .8. We have implementation of that input for:

- Neo4jInput — input source for working with graph database Neo4J
- GraphxInput — input source for working with graph presented in memory using GraphX library
- LinearInput — input source for working with linear input data like strings

```
trait Input[+L, +N] {
  def filterEdges(nodeId: Int,
      predicate: L => Boolean): Seq[(L, Int)]
  def checkNode(nodeId: Int,
      predicate: N => Boolean): Option[N]
}
```

Figure 8: Generalized input interface

3.4 Semantic Actions

4 EVALUATION

In this section, we present an evaluation of our graph querying library. We measure its performance on a classical set of ontology graphs [36]: both when the graph is loaded into the RAM and for the integration with the Neo4j database and compare it with the solution based on the GLL parsing algorithm [7] and the Trails [13] library for graph traversals. We also show how may-alias static code analysis can be done by means of the developed library.

All tests have been performed on a computer running Fedora 27 with quad-core Intel Core i7 2.5 GHz CPU and 8 GB of RAM.

4.1 Ontology querying

Querying for ontologies is a well-known graph querying problem. We evaluate our library on some popular ontologies which are presented as RDF files in the paper [36]. First, we convert RDF files to a labelled directed graph in the following manner: we create two edges (*subject*, *predicate*, *object*) and (*object*, *predicate^{-1}*, *subject*) for every RDF triple (*subject*, *predicate*, *object*). Then the graph is either loaded into the Neo4j database or is loaded directly into the memory. Then we run two queries from the paper [7] for these graphs. The grammars for the queries are presented in Fig. 10 and Fig. ??.

```
val query1 = syn(
    "subclassof -1" ~ S.? ~ "subclassof" |
    "type -1" ~ S.? ~ "type")
```

Figure 9: The same generation query (Query 1) in Meerkat

```
val S = syn(
  "subclassof -1" ~ S ~ "subclassof")
val query2 = syn(S ~"subclassof")
```

Figure 10: Query 2 grammar

The performance results are shown in the table 2. #*triples* reflects the size of the RDF file with the corresponding ontology, #*results* is a number of pairs of the nodes between which at least one S-path exists.

The Meerkat-based and the GLL-based [7] solutions show the same results (column #*results*) and the Query 1 runs up to two times faster on the Meerkat-based solution than on the GLL-based, meanwhile the GLL-based solution is faster for the Query 2. Querying the database is naturally $2 - 4$ times slower than querying the graphs located in the RAM.

4.2 Static code analysis

Alias analysis is a fundamental static analysis problem [17]: it checks may-alias relations between code expressions and can be formulated as a context-free language (CFL) reachability problem [25] which is closely related to context-free path querying problem. In this analysis, a program is represented as a Program

$$M \rightarrow \overline{D}\, V\, D$$
$$V \rightarrow (M?\, \overline{A})^*\, M?\, (A\, M?)^*$$

Figure 11: Context-Free grammar for the may-alias problem

Expression Graph (PEG) [37]. Vertices in a PEG correspond to program expressions and edges are relations between them. There are two types of edges possible while analyzing C source code: **D**-edge and **A**-edge.

- Pointer dereference edge (**D**-edge). For each pointer dereference $*e$ there is a directed **D**-edge from e to $*e$.
- Pointer assignment edge (**A**-edge). For each assignment $*e_1 = e_2$ there is a directed **A**-edge from e_2 to $*e_1$

For the sake of simplicity, we add edges labelled by $\overline{D}$ and $\overline{A}$ which correspond to the reversed **D**-edge and **A**-edge, respectively.

The grammar for may-alias problem from [37] is presented in Fig. 11. It contains two nonterminals **M** and **V** and allows us to make two kinds of queries for each of them.

- **M**-production means that two l-value expressions are memory aliases i.e. may refer to the same memory location.
- **V**-production means that two expressions are value aliases i.e. may evaluate to the same pointer value.

We run the **M** and **V** queries on some open-source C projects: the results are in table 3. We can conclude that our solution is expressive and performant enough to be used for static analysis problems which can be expressed as CFPQs.

```
val M = syn("nd" ~ V ~ "d")
val V = syn((M.? ~ "na").* ~ M.? ~ ("a" ~ M.?).*)
```

Figure 12: Meerkat representation of may-alias problem grammar

4.3 Comparison with Trails

Trails [13] is a Scala graph combinator library. It provides traversers for describing paths in graphs which resemble parser combinators and calculates possibly infinite stream of all possible paths described by the composition of basic traversers. Both Trails and Meerkat-based solution support parsing of the graphs located in RAM, so we compare the performance of Trails and Meerkat-based solution on the ontology queries described above: the results are presented in table 2. Trails and Meerkat-based solution compute the same results, but Trails is up to 10 times slower than Meerkat-based solution.

To summarize, we demonstrated that parser combinators are expressive enough to be used for implementing real queries. Our implementation is as performant as the other existing combinators library and is comparable to the GLL-based solution.

5 CONCLUSION

We propose a native way to integrate a language for context-free path querying into a general-purpose programming language. Our

Ontology	#tripples	#results	In memory graph (ms)	Query 1					Query 2				
				DB query (ms)	Trails (ms)	GLL (ms)	#results	In memory graph (ms)	DB query (ms)	Trails (ms)	GLL (ms)		
atom-primitive	425	291	685	15454	112	167	2849	232	122	49	52	453	19
biomedical-measure-primitive	459	341	711	15156	226	247	3715	482	2871	34	42	60	26
foaf	631	256	815	4118	16	25	432	29	10	1	2	1	1
funding	1086	778	1480	17634	123	152	367	179	1158	18	23	76	13
generations	273	129	351	2164	6	21	9	12	0	0	0	0	0
people_pets	640	337	834	9472	63	84	75	80	37	2	3	2	1
pizza	1980	671	2604	56195	544	650	7764	793	1262	44	47	905	50
skos	252	144	323	810	4	9	6	6	1	0	1	0	0
travel	277	131	397	2499	21	55	34	21	63	2	2	1	2
univ-bench	293	179	413	2540	15	43	31	24	81	2	2	2	1
wine	1839	733	2450	66572	543	727	3156	606	133	5	7	4	5

Table 2: Comparation of Meerkat, Trails and GLL performance on ontologies

Program	#edges	#nodes	Code Size (KLOC)	In memory graph (ms)	Neo4j graph (ms)	M-aliases	N-aliases
wc-5.0	332	770	0.5	0	4	174	107
pr-5.0	815	2062	1.7	13	17	131	63
ls-5.0	1687	4734	2.8	52	76	682	255
bzip2-1.0.6	632	1508	5.8	9	80	75	
gzip-1.8	2687	7510	31	120	182	4507	227

Table 3: Running may-alias queries on Meerkat on some C open-source projects

solution can handle arbitrary context-free grammars and arbitrary input graphs. The proposed approach is language-independent and may be implemented for closely all general-purpose programming languages. We implement it in the Scala programming language and show that our implementation can be applied to the real world problems.

We can propose some possible directions for the future work. First of all, it is necessary to formulate the creation of a user-friendly interface for SPPF processing. We can just extract reachability information, but SPPF contains mach more useful information. One such representation may be a set of paths with additional information about their structure. This may simplify debugging and query result processing.

In order to improve performance and investigate scalability of proposer solution it is necessary to try to implement parallel single machne and distributed GLL. It is not only algorithmic problem: to get practical solution we should choose appropriate tools, libraryes for parallel and distributet computing for Scala.

Another direction is a semantic actions computation, otherwise known as attributed grammars handling. It increases the expressiveness of queries by means of the specification of user-defined actions, such as filters, over subqueries result. Although it is impossible in general, techniques such as lazy evaluation can provide a technically adequate solution. Anoter possible direction is utilizatuin of relational programming (minikanren) which is aimed to search [?]. For what class of semantic actions it is possible to provide a precise general solution is a theoretical question to be answered.

Some important problems in static code analysis reqwuire languages more expressive than context-free one. For example, context-sensitive data-dependence analysis may be precisely expressed in terms of linear-conjunctive language [21] reachability, but not context-free [35]. While problem formulation is precise, it is possible to get only approximated solution, because emptyness problem for linear-conjunctive languages is undersidable. It would be an interesting task to support not only linear-cinjunctive grammars, but arbitrary conjunctive grammars [20] in te library and investigate nature of approximation.

All these improvements may provide !!! Declarative code analysis tools and languages Flix [16]. Reps [26]

Improved version of OpenCypher [15], which is the one of the most popular graph query languages, provides context-free path querying mechanism. Detailed comparison with it may provide more information for direcion of future work.

ACKNOWLEDGMENTS

REFERENCES

[1] Serge Abiteboul, Richard Hull, and Victor Vianu. 1995. Foundations of Databases. (1995).

[2] Pablo Barceló, Gaelle Fontaine, and Anthony Widjaja Lin. 2013. Expressive Path Queries on Graphs with Data. In *International Conference on Logic for Programming Artificial Intelligence and Reasoning*. Springer, 71–85.

[3] Chris Barrett, Riko Jacob, and Madhav Marathe. 2000. Formal-language-constrained path problems. *SIAM J. Comput.* 30, 3 (2000), 809–837.

[4] Osbert Bastani, Saswat Anand, and Alex Aiken. 2015. Specification inference using context-free language reachability. In *ACM SIGPLAN Notices*, Vol. 50. ACM, 553–566.

[5] Tevfik Bultan, Fang Yu, Muath Alkhalaf, and Abdulbaki Aydin. 2018. String Analysis for Software Verification and Security. (2018).

[6] James Cheney, Sam Lindley, and Philip Wadler. 2013. A Practical Theory of Language-integrated Query. *SIGPLAN Not.* 48, 9 (Sept. 2013), 403–416. https://doi.org/10.1145/2544174.2500586

[7] Semyon Grigorev and Anastasiya Ragozina. 2017. Context-free Path Querying with Structural Representation of Result. In *Proceedings of the 13th Central & Eastern European Software Engineering Conference in Russia (CEE-SECR '17)*.

ACM, New York, NY, USA, Article 10, 7 pages. https://doi.org/10.1145/3166094.3166104

[8] Jelle Hellings. 2014. Conjunctive context-free path queries. (2014).

[9] Jelle Hellings. 2015. Path Results for Context-free Grammar Queries on Graphs. *CoRR* abs/1502.02242 (2015). http://arxiv.org/abs/1502.02242

[10] Graham Hutton and Erik Meijer. 1996. Monadic parser combinators. (1996).

[11] Anastasia Izmaylova, Ali Afroozeh, and Tijs van der Storm. 2016. Practical, General Parser Combinators. In *Proceedings of the 2016 ACM SIGPLAN Workshop on Partial Evaluation and Program Manipulation (PEPM '16)*. ACM, New York, NY, USA, 1–12. https://doi.org/10.1145/2847538.2847539

[12] Mark Johnson. 1995. Memoization in Top-down Parsing. *Comput. Linguist.* 21, 3 (Sept. 1995), 405–417. http://dl.acm.org/citation.cfm?id=216261.216269

[13] Daniel Kröni and Raphael Schweizer. 2013. Parsing Graphs: Applying Parser Combinators to Graph Traversals. In *Proceedings of the 4th Workshop on Scala (SCALA '13)*. ACM, New York, NY, USA, Article 7, 4 pages. https://doi.org/10.1145/2489837.2489844

[14] Kazumasa Kumamoto, Toshiyuki Amagasa, and Hiroyuki Kitagawa. 2015. A System for Querying RDF Data Using LINQ. In *Network-Based Information Systems (NBiS), 2015 18th International Conference on*. IEEE, 452–457.

[15] Tobias Lindaaker. 2017. OpenCypher Path Patterns (CIP2017-02-06 Path Patterns). (2017). https://github.com/thobe/openCypher/blob/rpq/cip/1.accepted/CIP2017-02-06-Path-Patterns.adoc#153-compared-to-context-free-languages

[16] Magnus Madsen, Ming-Ho Yee, and Ondřej Lhoták. 2016. From Datalog to Flix: A Declarative Language for Fixed Points on Lattices. In *Proceedings of the 37th ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI '16)*. ACM, New York, NY, USA, 194–208. https://doi.org/10.1145/2908080.2908096

[17] Thomas J. Marlowe, William G. Landi, Barbara G. Ryder, Jong-Deok Choi, Michael G. Burke, and Paul Carini. 1993. Pointer-induced Aliasing: A Clarification. *SIGPLAN Not.* 28, 9 (Sept. 1993), 67–70. https://doi.org/10.1145/165364.165387

[18] Erik Meijer, Brian Beckman, and Gavin Bierman. 2006. LINQ: Reconciling Object, Relations and XML in the .NET Framework. In *Proceedings of the 2006 ACM SIGMOD International Conference on Management of Data (SIGMOD '06)*. ACM, New York, NY, USA, 706–706. https://doi.org/10.1145/1142473.1142552

[19] A. Mendelzon and P. Wood. 1995. Finding Regular Simple Paths in Graph Databases. *SIAM J. Computing* 24, 6 (1995), 1235–1258.

[20] Alexander Okhotin. 2001. Conjunctive grammars. *Journal of Automata, Languages and Combinatorics* 6, 4 (2001), 519–535.

[21] Alexander Okhotin. 2003. On the Closure Properties of Linear Conjunctive Languages. *Theor. Comput. Sci.* 299, 1 (April 2003), 663–685. https://doi.org/10.1016/S0304-3975(02)00543-1

[22] Polyvios Pratikakis, Jeffrey S Foster, and Michael Hicks. 2006. Existential label flow inference via CFL reachability. In *SAS*, Vol. 6. Springer, 88–106.

[23] Eric Prud, Andy Seaborne, et al. 2006. SPARQL query language for RDF. (2006).

[24] Joan Gerard Rekers. 1992. *Parser generation for interactive environments*. Ph.D. Dissertation. Universiteit van Amsterdam.

[25] Thomas Reps. 1997. Program Analysis via Graph Reachability. In *Proceedings of the 1997 International Symposium on Logic Programming (ILPS '97)*. MIT Press, Cambridge, MA, USA, 5–19. http://dl.acm.org/citation.cfm?id=271338.271343

[26] Thomas Reps, Susan Horwitz, and Mooly Sagiv. 1995. Precise Interprocedural Dataflow Analysis via Graph Reachability. In *Proceedings of the 22Nd ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages (POPL '95)*. ACM, New York, NY, USA, 49–61. https://doi.org/10.1145/199448.199462

[27] Juan L Reutter, Miguel Romero, and Moshe Y Vardi. 2015. Regular queries on graph databases. *Theory of Computing Systems* (2015), 1–53.

[28] Marko A Rodriguez. 2015. The gremlin graph traversal machine and language (invited talk). In *Proceedings of the 15th Symposium on Database Programming Languages*. ACM, 1–10.

[29] Elizabeth Scott and Adrian Johnstone. 2010. GLL parsing. *Electronic Notes in Theoretical Computer Science* 253, 7 (2010), 177–189.

[30] Elizabeth Scott, Adrian Johnstone, and Rob Economopoulos. 2007. BRNGLR: a cubic Tomita-style GLR parsing algorithm. *Acta informatica* 44, 6 (2007), 427–461.

[31] Petteri Sevon and Lauri Eronen. 2008. Subgraph queries by context-free grammars. *Journal of Integrative Bioinformatics* 5, 2 (2008), 100.

[32] Niklaus Wirth. 1996. Extended Backus-Naur Form (EBNF). *ISO/IEC* 14977 (1996), 2996.

[33] Dacong Yan, Guoqing Xu, and Atanas Rountev. 2011. Demand-driven Context-sensitive Alias Analysis for Java. In *Proceedings of the 2011 International Symposium on Software Testing and Analysis (ISSTA '11)*. ACM, New York, NY, USA, 155–165. https://doi.org/10.1145/2001420.2001440

[34] Mihalis Yannakakis. 1990. Graph-theoretic methods in database theory. In *Proceedings of the ninth ACM SIGACT-SIGMOD-SIGART symposium on Principles of database systems*. ACM, 230–242.

[35] Qirun Zhang and Zhendong Su. 2017. Context-sensitive Data-dependence Analysis via Linear Conjunctive Language Reachability. In *Proceedings of the 44th ACM SIGPLAN Symposium on Principles of Programming Languages (POPL 2017)*.

ACM, New York, NY, USA, 344–358. https://doi.org/10.1145/3009837.3009848

[36] Xiaowang Zhang, Zhiyong Feng, Xin Wang, Guozheng Rao, and Wenrui Wu. 2016. Context-free path queries on RDF graphs. In *International Semantic Web Conference*. Springer, 632–648.

[37] Xin Zheng and Radu Rugina. 2008. Demand-driven Alias Analysis for C. In *Proceedings of the 35th Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages (POPL '08)*. ACM, New York, NY, USA, 197–208. https://doi.org/10.1145/1328438.1328464